Runaway edges: a smoothness-penalty blind spot makes Kolmogorov–Arnold networks non-identifiable on correlated inputs

Anonymous authors
Paper under double-blind review

Abstract

Kolmogorov–Arnold networks (KANs) are promoted as interpretable because each edge carries a learned univariate function. We show that a widely used family of smoothness penalties leaves a blind spot that undermines this promise: second-difference and curvature penalties do not charge for affine components of an edge, and SiLU base weights are often unpenalised. When the inputs to a trained nonlinear layer are correlated, affine components on different edges can cancel, so training inflates the parameters along directions that neither the penalty nor the data can see, and independent runs do not agree. In pre-registered experiments with a known teacher, correlated inputs gave 0 of 16 converged models, including under a recently proposed curvature penalty, while a small penalty on the affine null space gave 16 of 16 certified minima and essentially exact recovery of the teacher. Along certified regularisation paths spanning five decades, the loss behaves as if its infimum lies at infinity, on synthetic data and on real electrolyte-conductivity data; closing the null space bounds it. Non-convergence rises with input correlation, and the runaway and its fix carry over from Gaussian-RBF to B-spline edges. On the real data, polishing beyond early stopping degraded predictions for a held-out co-solvent (RMSE 0.142 $\rightarrow$ 1.03 under a weak null-space penalty; 0.33 under a stronger one). Certified minima can still be degenerate or distinct, so convergence and identifiability are separate problems. We give practical recommendations for penalised KANs.

1 Introduction

KANs replace a network's fixed activations with learned univariate functions on every edge. Their appeal for science rests on reading those functions: an edge that looks like an exponential decay suggests a physical law. That reading presumes the edges are determined by the data. If different training runs produce different edges with equally good predictions, the edges describe the optimiser, not the system.

This paper identifies a concrete, common mechanism that breaks that presumption. It arises from the interaction of three ingredients that are each standard in practice:

1. **A smoothness penalty with an affine null space.** Second-difference (P-spline) penalties and curvature penalties ignore straight-line components of an edge; SiLU base terms are frequently unpenalised.
2. **Correlated inputs to a layer.** Physical descriptors often move together. Hidden units can also become correlated during training, exposing the next layer.
3. **Training toward convergence.** Early stopping hides the problem; longer optimisation exposes it.

Together they create directions along which edges grow into large, mutually cancelling affine shapes. The penalty does not see them, and on the training data their contributions cancel. The loss keeps decreasing

slightly along these directions, and over the range we could test it behaves as if it had no finite minimum. Seeds wander apart, no model certifies as converged, and on inputs off the training correlation pattern the cancellation fails.

Contributions.

- We describe the mechanism: a penalty's affine blind spot plus correlated inputs to a trained nonlinear layer. A layer whose output is linear in its parameters cannot run away (Sections 3 and 6).
- In a pre-registered synthetic study, a recently proposed curvature penalty does not prevent the runaway on correlated inputs, while a small penalty on the affine null space does and restores recovery of a known solution (Section 4).
- On real electrolyte-conductivity data, we trace a complete failure to converge to this mechanism and document its effect on a held-out co-solvent (Section 5).
- Along certified regularisation paths, the loss behaves as if its infimum lies at infinity over five decades of regularisation, on synthetic and real data; closing the null space bounds it (Section 6).
- Non-convergence rises with input correlation, and the runaway and its fix carry over from Gaussian-RBF to B-spline edges (Sections 7 and 8).
- Certified minima can be degenerate, distinct or not minima at all; we separate three kinds of flatness and give recommendations (Sections 9 and 10).

All experiments were pre-registered: predictions, thresholds and interpretation rules were committed before the code that tested them. Of 65 predictions in the real-data case study, 11 held, 22 failed and 32 were inconclusive; we report them as recorded. The revision studies of Sections 6–8 and the applicability analysis followed the same protocol; Appendix A lists their verdicts, and Appendix D every change made after an outcome was known.

2 Background and related work

KAN edges. A KAN layer maps inputs to outputs by summing learned univariate functions, one per edge. In the original formulation (Liu et al., 2025) each edge is a SiLU base term plus a B-spline expansion. FastKAN (Li, 2024) approximates the B-splines with Gaussian radial basis functions on fixed centres, which is the edge type used in our experiments:

$$\varphi(x) = \alpha\,\mathrm{SiLU}(x) + \sum_{k=1}^{K} c_k \exp\left(-\left(\frac{x-\mu_k}{h}\right)^2\right) \tag{1}$$

Smoothness penalties and their null spaces. P-spline penalties charge squared second differences of the coefficients (Eilers & Marx, 1996). Constant and linear coefficient patterns cost nothing: the penalty's null space is affine. In generalised additive models this is well understood, and Marra & Wood (2011) proposed an additional penalty on the null space (the "double penalty") for variable selection; Wood (2017) treats null spaces and concurvity in additive models in depth. Our fix is the same construction; what is new is its consequence inside a KAN, where edges are summed within units and passed through further layers.

Curvature penalties for KANs. Bagrow (2026) proposed per-edge curvature penalties for compositional smoothness, combining the second-difference penalty with a term $K_{\mathrm{silu}}\,\alpha^2$ on the SiLU weight, and noted that the penalty has affine functions as its null space by design. That work evaluated on independent, uniformly sampled inputs. Second-derivative penalties on spline coefficients also appear in other KAN variants, typically on the spline part only, leaving the SiLU base term free (Cang et al., 2024; see also the free-knots KAN of Zheng et al., 2025). We show that the affine null space becomes a runaway direction when inputs are correlated, and that penalising the SiLU weight alone does not close it.

Concurvity. In additive models, including neural additive models (Agarwal et al., 2021), correlated inputs let one input's shape function absorb another's, making shape functions non-identifiable; this is known as concurvity. Siems et al. (2023) proposed a regulariser penalising correlations between transformed features. The cross-edge cancellation we describe is concurvity inside a KAN unit. In an additive model with a penalised null space it shifts shape between functions; combined with an unpenalised affine null space in a KAN, it becomes unbounded growth.

Losses without a finite minimum. When a loss attains its infimum only at infinity, gradient descent drives the parameter norm to diverge, slowly, and optimisation can continue while held-out performance worsens (Soudry et al., 2018). Our runaway behaves like an instance of this, created by a penalty's null space. It differs from the benign flat directions of linear least squares, where gradient descent converges to the minimum-norm solution: here the direction is slightly tilted, so gradient descent keeps following it. On such ill-conditioned problems, early stopping acts much like ridge regularisation (Yao et al., 2007; Ali et al., 2019).

Parameter-space symmetries. Exact symmetries (unit permutations, sign flips) make parameters non-identifiable in tanh and RBF networks, and completeness results characterise them (Sussmann, 1992; Kůrková & Neruda, 1994). KANs add continuous symmetries such as affine reparametrisations of hidden units. The runaway is not an exact symmetry: it is an approximate one, a nearly flat but tilted direction, which is what prevents convergence.

3 Mechanism

Consider a unit that sums edges from two inputs, $s = \varphi_1(x_1) + \varphi_2(x_2)$, each edge as in Section 2, trained with a second-difference penalty on its RBF coefficients.

Step 1: a direction the penalty cannot see. Add a linear pattern $t\,\ell$ to the RBF coefficients of edge 1 and subtract it from edge 2, with ℓ_k proportional to the centre index. Linear coefficient patterns have zero second difference, so the penalty is unchanged for every t (Figure 1b). With Gaussian bumps on a regular grid, the pattern produces a nearly affine function: over the central 80% of the input range it deviates from a straight line by at most 2% of its range (Figure 1a), and by up to 13% near the grid edges. Edge 1 thus gains approximately $t\,a\,x_1$ and edge 2 loses $t\,a\,x_2$.

Step 2: a direction the data cannot see. The unit's pre-activation changes by approximately $t\,a\,(x_1 - x_2)$. If x_1 and x_2 are strongly correlated, $x_1 - x_2$ is small wherever the data lie, so predictions change little even for large t (Figure 1c). The direction is nearly free in both the penalty and the data.

Step 3: a receding minimum. "Nearly free" is not "free". Moving along the direction lets the model shed curvature into the free affine and SiLU components, so the loss keeps decreasing slightly. Both terms contribute (post-hoc decomposition): in the synthetic runaway models the smoothness penalty fell to about a quarter of its early value and supplied a median 61% of the decrease while training error fell only 5%; on real data, from early stopping to the capped state, the penalty fell 8–88× and supplied 17–36% of the decrease, the data term the rest. The growth went into whatever the penalty left free: the linear part of the RBF coefficients grew 34–40× in the synthetic runaway cells, and the SiLU weights 85× where they were unpenalised. In 12 of 20 synthetic runaway models, training error fell while test error rose. Over the range we could test, the loss behaves as if its infimum lies at infinity: along a certified L2 regularisation path, the minimisers grow without levelling off over five decades of the regularisation weight while the loss keeps falling, and they recede along a curved valley rather than a straight ray (Section 6). Growth alone is not the signature, because a distant finite minimum also produces large growth (Step 4); what distinguishes the runaway is that the optimisation never reaches a certified minimum. Different seeds drift along different such directions.

Step 4: the runaway needs a trained nonlinear layer. If the parameters being trained enter the output linearly, as in the output layer when the hidden activations are held fixed, or in a one-layer KAN, the penalised loss is a convex quadratic. It may have flat directions, but no tilted ones, so it cannot run away: with the first layer frozen, every model reached a certified finite minimum, and so did a one-layer model on correlated inputs, although the minimum could lie far away (growth up to ×583 on real data). The runaway

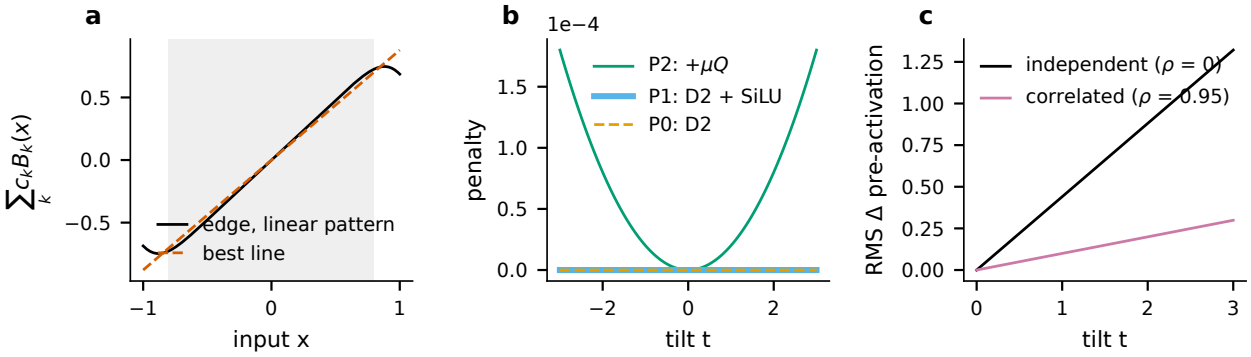


Figure 1: Mechanism (computed illustration) for the edge type of Section 2 (8 Gaussian RBFs on $[-1, 1]$). **(a)** Edge function produced by a unit linear coefficient pattern and its best straight line; the shaded band is the central 80% of the range, where the maximum deviation is 2.0% of the function's range (12.9% over the full range). **(b)** Penalty as a function of a tilt t applied as $+t$ to edge 1 and $-t$ to edge 2: the second-difference penalty (P0) and P1 are exactly zero for every t; only P2's term μQ grows ($\lambda = 10^{-3}$, $\mu = 10^{-5}$). **(c)** RMS change in the unit's pre-activation per unit tilt: 0.0996 on correlated inputs ($\rho = 0.95$) against 0.4403 on independent inputs (4.42×).

therefore needs trained parameters that act through a nonlinearity. Correlation can reach such a layer from correlated inputs, or from hidden units that become correlated during training, so the output layer can host the runaway even when the raw inputs are independent, as long as the first layer also adapts. Correlated hidden units are necessary but not sufficient: they also occur in models that converge once the null space is penalised. Whether a given nonlinear problem degenerates depends on the data: with the second layer frozen, the synthetic first-layer problems settled cleanly, while the real first-layer problem, with few-valued, strongly correlated descriptors, reached only highly degenerate minima, one of which was not a minimum at all.

Consequences.

- **Non-identifiability:** edges differ across seeds by large affine components that leave predictions almost unchanged; even certified minima can be degenerate or distinct from one another.
- **Non-convergence:** second-order optimisers stall and certification fails. Since growth is also large on the way to a distant finite minimum, certification, not growth, is the criterion.
- **Extrapolation risk:** on inputs outside the training data's support, the inflated components need not cancel. In our real-data case, polishing damaged predictions for a held-out co-solvent lying outside the training descriptor hull.

The fix. Penalise the null space directly. For each edge, let P_{null} project its RBF coefficients onto constant-plus-linear patterns in the centre index, and add

$$Q = \sum_e \| P_{\text{null}}\, c_e \|^2 \tag{2}$$

with a small weight μ, in addition to a penalty on the SiLU weight. This makes every affine component cost something, so the regularised minimisers stay bounded. Penalising only part of the free directions is not enough: the growth moves into whatever stays unpenalised, such as biases or embeddings, so these need a small penalty too.

Table 1: Synthetic 2×3 study, per cell.

Cell	Norm growth	Certified	Unit match to reference
Independent, P0	12.5×	6/8	1.000
Independent, P1	7.9×	6/8	1.000
Independent, P2	1.09×	8/8	1.000
Correlated, P0	65.7×	0/8	0.774
Correlated, P1	32.1×	0/8	0.883
Correlated, P2	1.00×	8/8	1.000

4 Controlled study

Design. A teacher KAN [4, 3, 1] with smooth, hand-specified edges (fitted by smoothness-penalised least squares) generates targets with Gaussian noise of 5% of the output spread; n = 2000 training rows. Inputs are $\mathcal{N}(0, 1)/3$, clipped to the grid, either independent or correlated in pairs ($\rho = 0.95$ for x_1/x_2 and x_3/x_4). Students share the teacher's architecture and are trained under three penalty levels:

- **P0:** second-difference penalty on RBF coefficients ($\lambda = 10^{-3}$).
- **P1:** P0 plus $K_{\text{silu}}\,\alpha^2$ on each SiLU weight (Bagrow-style curvature penalty, $K_{\text{silu}} = 0.443$).
- **P2:** P1 plus μQ, the affine null-space penalty.

Each of the 6 cells has 8 seeds. Each student is trained with AdamW and early stopping, then polished with an exact-Hessian trust-region method and certified by a second-order criterion (gradient, curvature and Newton decrement). μ was calibrated before the experiment by a pre-registered rule: 10^{-6} failed the pilot (5.0× growth) and 10^{-5} passed (1.32×) and was frozen.

The teacher itself runs away. Polished under P0 and P1 on correlated inputs, the teacher's parameters grew 85× and 1069× without certifying, while its predictions stayed within half a noise standard deviation of the original. The runaway changes the parameters, not the function.

Medians over 8 seeds; growth is the squared parameter norm, polished over early-stopped.

Results.

- **The published-style penalty does not prevent the runaway.** Penalising the SiLU weight (P1) halved the growth on correlated inputs but left 0 of 8 models certified (Appendix E, Figure 6).
- **Closing the null space prevents it everywhere.** P2 gave 16 of 16 certified models with no growth, in both input conditions.
- **Without the runaway, the internals are recoverable.** Every cell without a runaway matched the reference's units and edges exactly (1.000), with seed-to-seed differences around 10^{-4}.
- **Independent inputs are not fully safe.** Under P0 and P1, most independent seeds grew substantially (8–12×) but settled at certified minima; 4 of 16 (seeds 0 and 4 in each cell) never certified.

Where the runaway lives. A post-hoc anatomy of the parameter change shows that in the correlated cells almost all growth sits in the output layer: the layer-1 share of the change was 0.110 (P0) and 0.024 (P1) (Appendix E, Figure 7a). It went into the affine part of the output edges' RBF coefficients and, under P0, into their SiLU weights. Cancellation across the edges entering the output rose from about 1.9 to 3.1–3.2, and no hidden unit saturated. The two uncertified independent-P0 seeds put 99.4% of their change into the output layer. This is Step 4 of Section 3. Among independent-input models, those that ran away were exactly those whose hidden units had become correlated (largest pairwise correlation 0.70, against 0.05 for

certified models; Appendix E, Figure 7b), while certified P2 models kept correlations up to 0.88 without running away (post-hoc).

Verdicts. Of six pre-registered predictions, one held (the certification pattern, X2) and five failed. The failures refine rather than contradict the mechanism. Independent cells exceeded the growth threshold, mostly through finite growth. Cancellation and extrapolation effects were weaker than predicted. Recoverability failed only because the runaway cells were more recoverable than predicted, not less. In this synthetic setting, polishing did not degrade extrapolation: most of the off-pattern error in correlated cells was present at early stopping.

5 Real-data case study: electrolyte conductivity

Task and model. The CALiSol-23 challenge (de Blasio et al., 2024) asks for the log conductivity of lithium-ion electrolytes; every test sample contains diethyl carbonate (DEC), which never appears in training. KAN-δ anchors each prediction on the measured conductivity of pure propylene carbonate (PC) and learns only the correction, as in Δ-machine learning (Ramakrishnan et al., 2015), $\log k = \log k_{\mathrm{PC}} + x_{\mathrm{co}}\, h(z)$, where h is a FastKAN-style network of shape [10, 6, 1]. Its inputs are physical descriptors of the solvents (dielectric constant, log viscosity, molar mass), temperature, molality and a salt embedding. A second-difference penalty acts on the RBF coefficients, and the SiLU weights are unpenalised. Early-stopped after 107 AdamW epochs, it reached RMSE 0.142 on a held-out co-solvent (fold A: ethyl acetate held out). Its AdamW weight decay played no part: retraining without it left the early-stopped norms essentially unchanged and fold A's RMSE at 0.1420 (a two-sided pre-registered prediction whose verdict was the branch "incidental").

The inputs are correlated and few-valued. The three co-solvent descriptors take only 30 distinct combinations in training, and dielectric constant and log viscosity correlate at r = 0.966.

Nothing ever converged.

- Of 190 pairs of polished independent seeds, none computed the same function.
- After 20,000 float64 L-BFGS iterations (the "capped" state used below), 0 of 26 models converged; exact-Hessian trust-region methods also failed on every seed.
- Students trained on a frozen teacher matched its predictions (correlation 0.9988) but not its hidden units (0.66).

The anatomy matches Section 3.

- From early stopping to the capped state, the squared parameter norm grew 150–1550× while the smoothness penalty fell 8–88×.
- At least 99.9% of the coefficient change during exact-Newton polishing lay in the penalty's affine null space, in 9 of 9 models.
- Cancellation within an edge (SiLU against RBF) stayed small, 1.13–1.33. Cancellation across the edges entering a unit rose from 1.7–2.9 to 4–10, consistent with the correlated descriptors.
- The runaway spread across both layers: 30–40% of the parameter change went into layer-1 affine RBF coefficients, 37–38% into SiLU weights and 19–31% into layer-2 affine coefficients (post-hoc), and 12–39% of hidden units saturated. In the synthetic runaway the output layer dominated.

Extrapolation suffered. Polishing beyond early stopping under a weak null-space penalty ($\mu = 10^{-6}$) raised fold A's RMSE from 0.142 to 1.03 (0.33 under the stronger penalty $\mu = 10^{-4}$; Section 9). Fold A's co-solvent lies outside the training data's descriptor hull (Appendix E, Figure 8c), although on the correlation pattern of dielectric constant and viscosity. Within the fold, the damage scales with the co-solvent fraction: the learned correction enters as $x_{\mathrm{co}}\, h$, so any change polishing makes to h is magnified where x_{co} is large

(descriptive; Appendix E, Figure 8a, b). Across the co-solvent folds with at least 20 test rows, no distance measure distinguishably predicted where errors were large, and a planned predictive follow-up was dropped because its pre-stated precondition was not met. The degradation was much stronger than in the synthetic study, where correlation itself caused most of the off-pattern error (Section 11).

6 Is the minimum at infinity?

Failing to converge does not by itself show that the loss has no finite minimum. The minimum could be finite but very far away, or the problem could simply be too ill-conditioned for the optimiser to finish. Two pre-registered experiments separate these possibilities: freezing one layer, and following a certified regularisation path.

Freezing a layer. With the first layer frozen at its early-stopped weights, only the output layer is trained. Its parameters enter the output linearly, so the penalised loss is a convex quadratic. Every model then reached a certified minimum: 6 of 6 real KAN-δ seeds and 8 of 8 in each correlated synthetic cell, with no negative curvature at any logged iterate; a one-layer KAN trained on correlated inputs also certified, 8 of 8. The pre-registered prediction that the norm would grow at most threefold failed: the real models grew a median $\times 583$ to reach their minima. That growth is not a runaway, because it ends at a certified point. Freezing the second layer instead, in an exploratory arm, left the synthetic models certifying but the real first layer reaching only highly degenerate endpoints (Appendix F.1). Section 3 draws the conclusion: the runaway needs trained parameters acting through a nonlinearity, and whether such a problem degenerates depends on the data.

A certified regularisation path. For a loss L that is bounded below and a penalty R that grows without bound with the parameters, the global minimisers of $L + \gamma R$ stay bounded as $\gamma \to 0$ if and only if L has a finite minimiser. We use $\mathrm{R} = \|\theta\|^2$ over all parameters, follow γ from 10^{-2} down to 10^{-7} in half-decade steps, warm-start each value from the previous solution, and certify every point; degenerate points are accepted only after probes for remaining descent. On synthetic data, with the blind spot open, the minimisers' squared norm grew as a power of $1/\gamma$ (exponent 0.63–0.76 under P0, 0.36 under P1) while the unregularised loss kept falling; with the null space closed (P2), the path levelled off (exponent 0.014, ratio 1.004 between the two smallest γ). On the six real models, the same divergence appeared on six distinct branches: exponents 0.74–0.88, squared norms reaching 596–3052, and the unregularised loss falling from $2.05{\cdot}10^{-2}$ to 8.7–$9.3{\cdot}10^{-3}$. All pre-registered predictions held, the real-data discrimination check being taken from the synthetic controls as pre-registered. (Figure 2)

Two details matter for how far these replicate. On synthetic data the warm-started paths start from nearly one function; under P1 and P2 they stay together along the whole path, while under P0 they separate into a few branches. The real paths, by contrast, are genuinely distinct and compute different functions: at $\gamma \leq 10^{-4}$ their predictions differ by a median 13.7–16.3% of the output's standard deviation.

Robustness (post hoc). Polishing from independent cold starts found no small-norm basin at any γ; at $\gamma = 10^{-7}$ under P1, 8 of 9 cold starts reached a *lower* regularised objective than the warm path, at a larger norm. The warm path therefore understates the divergence. A straight line along the final drift direction rises almost at once: the minimisers recede along a curved valley, not along a ray. The growth sits in the penalty's blind directions, the affine RBF coefficients and SiLU weights of both layers, while the curved coefficients and the salt embedding shrink.

Closing the null space on real data. In a control whose two possible outcomes were fixed in advance, the same path on the real loss with the null-space penalty at $\mu = 10^{-4}$ levelled off: squared norms of 27.7 and 27.4 against 596 and 2200 for the same seeds with the blind spot open (ratios 1.09 and 1.10). The plateau is not a full stop: the norm still creeps up (whole-path exponents 0.105 and 0.124), in the biases, the curved RBF coefficients and the salt embedding, which this penalty leaves weakly penalised or free. The price is a higher training loss: +6% and +23% for the two seeds ($1.0926{\cdot}10^{-2}$ against $1.0300{\cdot}10^{-2}$, and $1.1224{\cdot}10^{-2}$ against $9.113{\cdot}10^{-3}$). When the penalty is instead lowered towards zero, the growth floods whichever parameters it leaves free: with the null-space penalty of Section 3 the salt embedding grows 44–218-fold, while a variant that also penalises biases and embeddings keeps the embedding bounded.

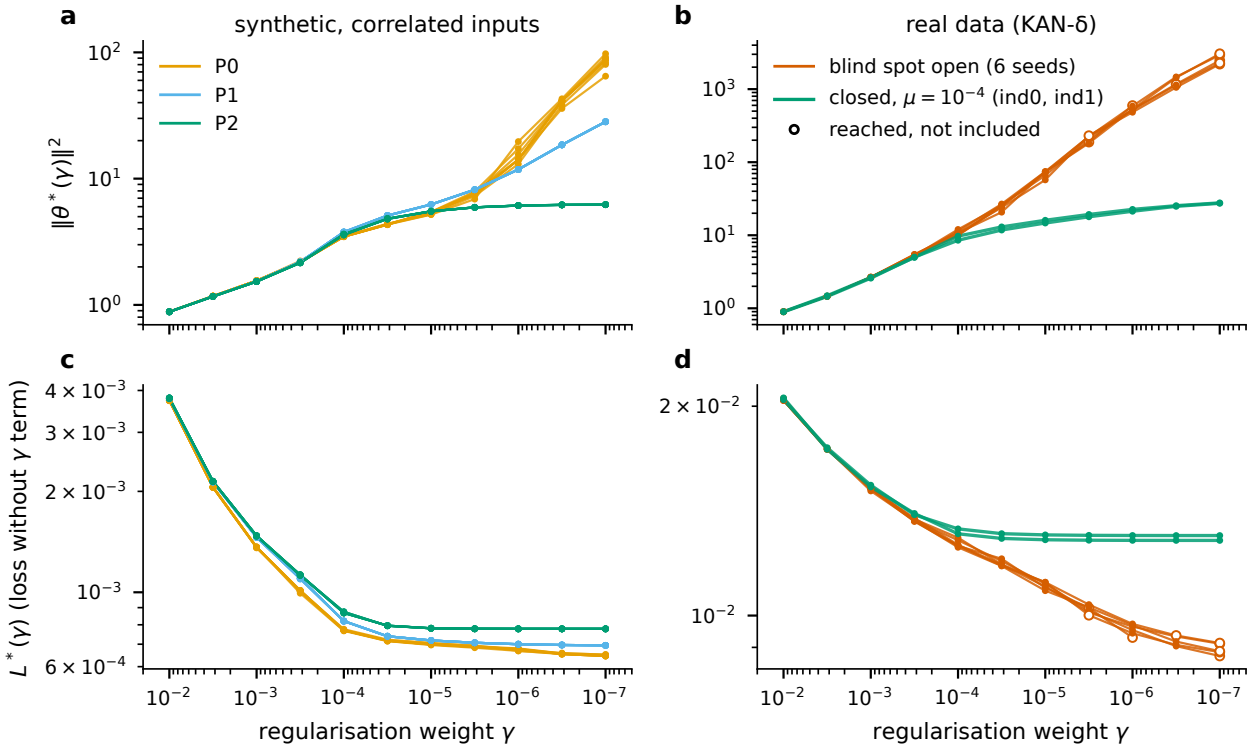


Figure 2: Certified L_2 regularisation paths, $\gamma = 10^{-2}$ to 10^{-7} in half-decade steps (warm-started), one line per seed. **(a, c)** Synthetic correlated inputs, 8 seeds per penalty level: $\|\theta^*\|^2$ grows as a power of $1/\gamma$ with exponent 0.63–0.76 (P0) and 0.36 (P1), and 0.014 with the null space closed (P2). **(b, d)** Real data: the 6 KAN-δ seeds with the blind spot open (exponents 0.74–0.88) and seeds ind0 and ind1 with the null-space penalty at $\mu = 10^{-4}$ (post-hoc control, outcomes fixed in advance; ratio between the two smallest γ 1.09 and 1.10). L^* is the loss without the γ term (for the closed paths, including μQ). Filled markers: points included in the pre-registered fits; open markers: reached but not included; absent: not reached.

Reading. Over the tested range, $\gamma = 10^{-2}$ to 10^{-7}, the loss with the blind spot open behaves as if its infimum lies at infinity, on synthetic and real data alike, and closing the null space bounds it. Two limits apply: the criterion concerns global minimisers while the path follows one local branch, and a finite minimum lying beyond the tested range cannot be excluded.

7 How much correlation does the blind spot need?

Section 4 compared two doses, independent inputs and $\rho = 0.95$. A pre-registered dose-response study fills in between: $\rho \in \{0, 0.5, 0.8, 0.9, 0.95, 0.99\}$ for both input pairs, P0 and P1 at every level, 8 seeds each, with paired random draws so that only the correlation changes, plus the null-space control P2 at $\rho = 0.99$. A replication gate ran first: the new $\rho = 0.95$ cells reproduced the committed synthetic study bit for bit.

Correlation drives non-convergence. The share of models that fail to certify rises with ρ (prediction held; P0 with one allowed inversion). It first exceeds one half at $\rho = 0.5$ under P0 and at $\rho = 0.9$ under P1, and from $\rho = 0.95$ no model certifies under either penalty. These thresholds belong to this synthetic setting ([4, 3, 1], Gaussian inputs); they are not thresholds for real data. With the null space closed, all 8 models certify even at $\rho = 0.99$, and a short regularisation path levels off (median D = 1.13, where D is the ratio of the squared norms at $\gamma = 10^{-7}$ and 10^{-5} along a short regularisation path). Along with the dose, the hidden units become correlated (largest pairwise correlation from about 0.05 at $\rho = 0$ to 0.82–0.96 at $\rho \geq 0.9$), and the polished training loss falls (Spearman -0.94): more correlation lets the runaway buy more loss. (Figure 3)

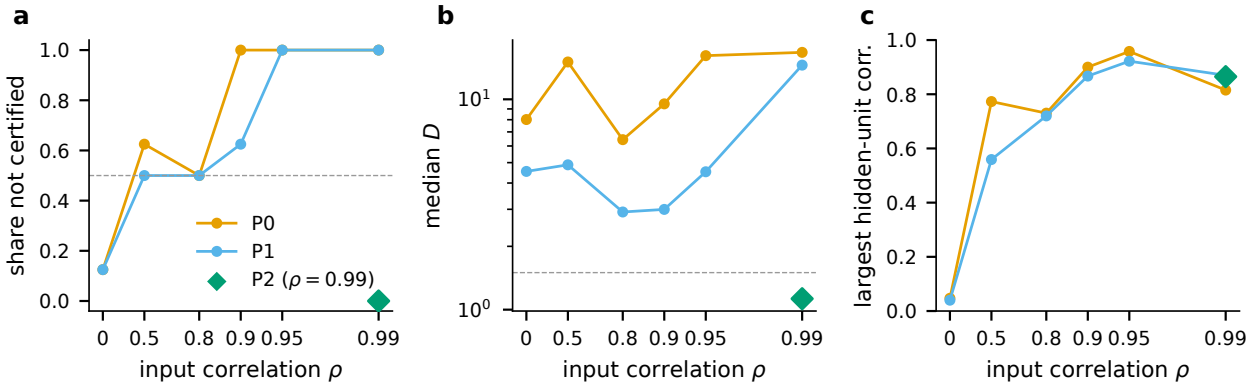


Figure 3: Correlation dose-response (synthetic [4, 3, 1], Gaussian inputs; 8 seeds per cell; the thresholds are synthetic, not real-data thresholds). Horizontal axis: ρ on a $\log 1/(1-\rho)$ scale. **(a)** Share of models not certified (dashed: one half). **(b)** Median short-path $D = \|\theta^*(10^{-7})\|^2/\|\theta^*(10^{-5})\|^2$ (log scale; dashed: 1.5). **(c)** Largest hidden-unit correlation, cell median. Diamonds: P2 at $\rho = 0.99$, 8 of 8 certified, median $D = 1.13$.

What a short path cannot tell. We also predicted that a short regularisation path ($\gamma = 10^{-3}$, 10^{-5}, 10^{-7}) would diverge more strongly as ρ grows. It did not: the path statistic D was already high at every dose, including $\rho = 0$. A post-hoc check with a rule fixed in advance showed why: where the polish certified, mostly at low ρ, 30 of 32 paths were heading towards that certified finite minimum, and a two-decade path cannot separate a distant finite minimum from a receding one (derivation in Appendix F.4). The receding minimum of Section 6 is therefore established at $\rho = 0.95$ and on the real data; at low ρ the minima are finite but distant; at intermediate ρ these tests do not settle the question. A secondary prediction, that the affine share of the parameter change would rise steadily with ρ, also failed.

8 Is it the basis?

The experiments so far use Gaussian-RBF edges. To test whether the runaway belongs to the basis or to the penalty's blind spot, we replicated the synthetic 2×3 study with cubic B-spline edges as in pykan and efficient-kan: uniform knots with G = 5 intervals on $[-1, 1]$, extended by three knots at each end, giving 8 basis functions per edge (matching the 8 Gaussians), plus the SiLU term. On these knots, constant and linear coefficient patterns reproduce constant and linear functions exactly, so the second-difference penalty's blind spot is exactly the affine functions. The teacher was refitted into the B-spline basis; inputs, seeds and noise were those of the RBF study.

The runaway carries over. With correlated inputs, none of the 16 P0 and P1 models certified (prediction held; Figure 4a). The runaway is a property of the penalty's affine null space, not of the Gaussian basis.

The fix carries over, with a lesson about calibration. The null-space weight was set by the same committed rule as in the RBF study (the smallest μ whose pilot models grew at most threefold), which here chose $\mu = 10^{-6}$, the bottom of the grid, against 10^{-5} for RBF edges. At that weight, the closed models certified 8 of 8 in both conditions, but the pre-registered bound on the short path failed with correlated inputs (median D = 2.23 against 1.5). Post-hoc checks, with predictions fixed in advance, traced this to the measurement window: the null-space penalty's curvature, about 1.8μ, lies between the two γ values that D compares. Extended to $\gamma = 10^{-9}$, the paths were bounded (Figure 4b, c), ending near the certified minima's norm though not always at the same point; and at the RBF study's $\mu = 10^{-5}$ the fix met the original standard (8 of 8 certified, D = 1.17 and 1.31). The growth-based rule chose a weight too small for a criterion carried over from another setting: a lesson about calibration, not about the fix.

Independent inputs. Fewer independent models certified than with RBF edges (5 and 4 of 8 under P0 and P1, against 6), so the pre-registered prediction failed, and its fixed reading applies: the B-spline basis

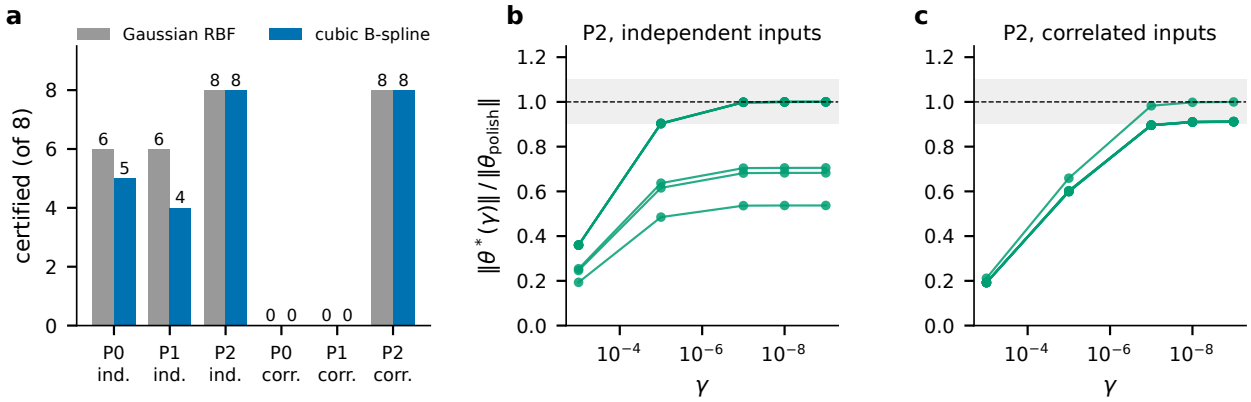


Figure 4: B-spline replication. **(a)** Certified models (of 8) per cell: the Gaussian-RBF study (frozen $\mu = 10^{-5}$) and the cubic B-spline replication (frozen $\mu = 10^{-6}$). **(b, c)** P2 regularisation paths at $\mu = 10^{-6}$, continued to $\gamma = 10^{-9}$ (post hoc, predictions fixed in advance): $\|\theta^*(\gamma)\|$ relative to each seed's certified polish minimum; shaded: the pre-fixed band [0.9, 1.1]. Median ratio at $\gamma = 10^{-9}$: 1.000 (independent) and 0.911 (correlated).

Table 2: Dose-response in μ on real data (seeds).

μ	Seeds at a numerical minimum	Negative curvature (seeds)	Norm growth (seeds)	Training fit vs runaway
10^{-6}	0 of 6	up to 9 eigenvalues	21–31×	−3% to +13%
10^{-5}	4 of 6	up to 4 eigenvalues	9–43×	12–22% worse
10^{-4}	6 of 6	none	1.6–16×	15–39% worse

is less benign on independent inputs. Post hoc, two of the seven failures look like runaways, one like slow convergence, and four are mixed; with 8 models per cell, the margin is one or two models.

A symmetry the RBFs lack. B-spline edges carry an exact 15-dimensional symmetry, constants moved between a unit's incoming edges and its bias, which holds only approximately for RBF edges; we excluded these analytically constructed directions from certification and from the degeneracy count (Appendix F.5). This was one of five changes to the study's calibration step made after outcomes were known; none touched a prediction, a threshold or an interpretation rule, and all are listed in Appendix D.

9 Closing the blind spot, and what remains

The null-space weight on real data. Adding μQ (with the SiLU weights included in Q) and polishing from the early-stopped weights:

Under a second-order certificate, all six seeds at $\mu = 10^{-4}$ were certified (Appendix E, Figure 9). (The originally pre-registered criterion, which required 10 stationary iterations, proved unreachable because the optimiser stops a few iterations after reaching a minimum; the certificate was defined afterwards and is reported alongside the original verdict.) Saturation fell to at most 0.5% and cross-edge cancellation to 2.2–2.9.

Against a regularised teacher, students recover its units and edges well, and the null-space penalty partly acts as gauge fixing (post hoc; Appendix F.3).

Convergence is not uniqueness. Independent seeds at certified minima still disagreed: their median prediction difference fell by only 26% (Figure 5b). With the runaway removed, the loss has several distinct isolated minima.

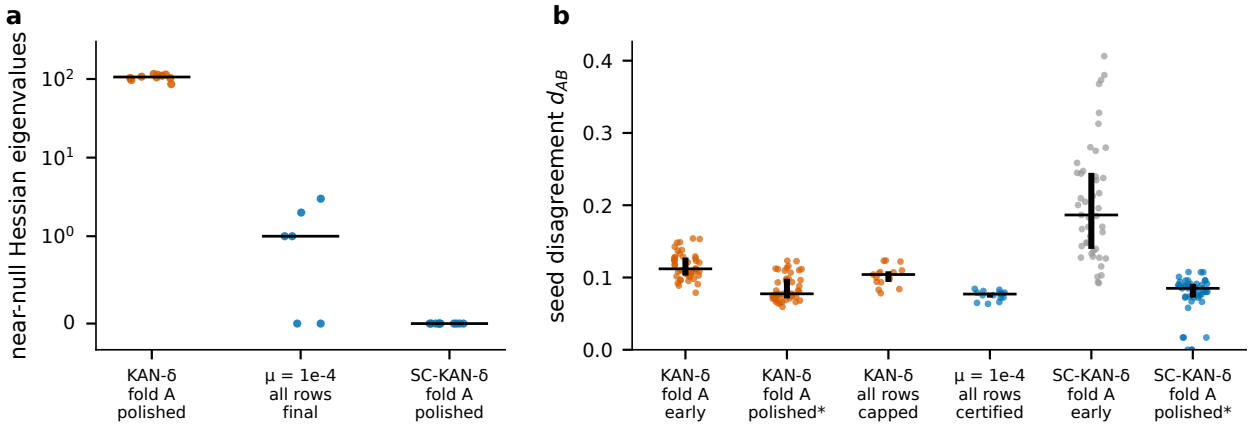


Figure 5: Convergence is not uniqueness. **(a)** Near-null Hessian eigenvalues ($|\lambda| < 10^{-8}\lambda_{\max}$) at the end of polishing: KAN-δ baseline on fold A, median 106 (range 86–116); null-space penalty $\mu = 10^{-4}$ on all rows, median 1 (range 0–3); SC-KAN-δ on fold A, 0 in all 10 models. **The baseline counts are taken at non-stationary points**: none of the 10 baseline models was certified, and all stopped at the 500-iteration cap. **(b)** Seed disagreement d_{AB} (median, interquartile range): KAN-δ fold A early 0.112; KAN-δ fold A polished* 0.077; KAN-δ all rows capped 0.104; $\mu = 10^{-4}$ all rows certified 0.077; SC-KAN-δ fold A early 0.187; SC-KAN-δ fold A polished* 0.085. Groups are evaluated on different row sets (fold A training rows or all rows), so values are not strictly comparable across row sets. Groups marked * (polished fold A) are not in any committed result: they were evaluated post hoc from the archived polished weights (SHA-verified; predictions on the 1434 fold A training rows), without any re-optimisation.

Three kinds of flatness. Certification needs care at flat points. Exact symmetries of the parameterisation, such as the constant shifts of B-spline edges (Section 8), are flat by construction and harmless; we exclude them analytically. Near-flat directions from poor conditioning are not harmless: certified points with many of them can be degenerate, and in one real-data case a point that passed the certificate kept descending on a longer restart while its squared norm grew 6.7-fold. The runaway valley itself is approximate and data-dependent. We therefore report the number of near-flat directions with every certificate, and probe degenerate points for remaining descent before counting them as minima. Certified, in short, is not the same as unique, and not necessarily the same as good: some certified minima lie at higher loss than points the regularisation path finds at smaller norm.

Extrapolation. On the held-out co-solvent, RMSE was 0.142 for the early-stopped model, 1.026 after polishing at $\mu = 10^{-6}$ and 0.33 at $\mu = 10^{-4}$. Bounding the runaway recovers most of the damage, but early stopping remains the most accurate choice: on an ill-conditioned problem, stopping early keeps the poorly determined directions near their starting values, much as ridge regularisation would (Yao et al., 2007; Ali et al., 2019).

Structured alternatives. Two architectures designed to remove runaway routes by construction, a shared-curve KAN and a two-anchor excess model, were less accurate than the early-stopped baseline under leave-one-co-solvent-out validation; the shared-curve model converged every time but its seeds still disagreed (Appendix F.2).

10 Recommendations for practitioners

1. **Penalise the blind spot, and everything else that is free.** If a KAN uses a second-difference or curvature penalty, add a small penalty on the constant-plus-linear component of each edge's coefficients and on each SiLU weight, in every layer, and give biases and embeddings a small penalty too: penalising only part of the free directions redirects the growth into the rest. This restores bounded minima, not the accuracy of early stopping.

2. **Calibrate that penalty by certification and a plateau, not by growth.** Choose its weight so that models certify and the regularisation path levels off; a rule based on norm growth chose a weight too small in our B-spline replication.
3. **Check input correlations, including between hidden units.** Strongly correlated or few-valued inputs make a trained nonlinear layer vulnerable, and hidden units that become correlated expose the next layer.
4. **Judge convergence by certification, not by norms or the loss.** Large growth can end at a finite minimum; a runaway is growth that never reaches a certified one. Report the number of near-flat directions with every certificate, and probe degenerate points before calling them minima.
5. **Do not read edges from unconverged models.** Compare edges across seeds only after removing exact symmetries, and only for models at certified minima.
6. **Treat seed disagreement as information.** Even at certified minima, distinct solutions remain; disagreement then signals that the data do not determine the function, not merely that training was unfinished.
7. **Validate on genuinely new conditions, and keep early stopping as a baseline.** A held-out fold from the training distribution can understate extrapolation error severalfold; in our real-data case, early stopping gave the best held-out accuracy of every variant tested.

11 Limitations and research process

Scope.

- One real dataset and one synthetic teacher; small networks ([4, 3, 1] and [10, 6, 1]).
- Two edge types: Gaussian RBFs on fixed centres and cubic B-splines on a fixed uniform grid, each with a SiLU base. Grid updates and other penalty families were not tested.
- Our P1 keeps the second-difference penalty and adds Bagrow's SiLU term; it is Bagrow-style, not a replication of his RBF curvature form.
- The receding minimum is established over $\gamma = 10^{-2}$ to 10^{-7} only, at $\rho = 0.95$ and on the real data; the dose thresholds are synthetic.
- The null-space weight in the shared-curve model was carried over from the synthetic calibration, not re-calibrated.

Post-hoc elements. The second-order certificate, the regularised-teacher comparison, the gauge-curvature analysis, the anatomy studies, the loss decomposition, the degeneracy probes, the cold starts, the dose-response settling check and the B-spline follow-up checks were defined after seeing results. We report them as exploratory and never in place of pre-registered verdicts; where their rules were fixed before computing, we say so.

Amendments and dropped work. Five changes to the B-spline study's calibration step were made after outcomes were known (Appendix D); none changed a prediction, a threshold or an interpretation rule. A planned predictive follow-up to the applicability analysis was dropped because its pre-stated precondition was not met.

Extrapolation. The strong degradation from polishing appeared on real data but not in the synthetic study, so we do not claim it as a general law.

Process. Each study committed its predictions, thresholds and interpretation rules before its code existed, and every verdict is linked to its commit; following up failed predictions, rather than reinterpreting them, led to the mechanism reported here. Appendix D describes the protocol, including the AI assistance used under it.

References

Rishabh Agarwal, Levi Melnick, Nicholas Frosst, Xuezhou Zhang, Ben Lengerich, Rich Caruana, and Geoffrey E. Hinton. Neural additive models: Interpretable machine learning with neural nets. In *Advances in Neural Information Processing Systems 34 (NeurIPS 2021)*, 2021. URL https://proceedings.neurips.cc/paper/2021/hash/251bd0442dfcc53b5a761e050f8022b8-Abstract.html.

Alnur Ali, J. Zico Kolter, and Ryan J. Tibshirani. A continuous-time view of early stopping for least squares regression. In *Proceedings of the 22nd International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 89 of *Proceedings of Machine Learning Research*, pp. 1370–1378. PMLR, 2019. URL https://proceedings.mlr.press/v89/ali19a.html.

James Bagrow. KANs need curvature: penalties for compositional smoothness. arXiv:2605.02190, 2026.

Yueyang Cang, Yu hang Liu, and Li Shi. Can KAN work? Exploring the potential of Kolmogorov–Arnold networks in computer vision. arXiv:2411.06727, 2024.

Paolo de Blasio, Jonas Elsborg, Tejs Vegge, Eibar Flores, and Arghya Bhowmik. CALiSol-23: Experimental electrolyte conductivity data for various Li-salts and solvent combinations. *Scientific Data*, 11:750, 2024. doi: 10.1038/s41597-024-03575-8.

Paul H. C. Eilers and Brian D. Marx. Flexible smoothing with B-splines and penalties. *Statistical Science*, 11(2):89–121, 1996. doi: 10.1214/ss/1038425655.

Věra Kůrková and Roman Neruda. Uniqueness of functional representations by Gaussian basis function networks. In *ICANN '94: Proceedings of the International Conference on Artificial Neural Networks*, pp. 471–474. Springer London, 1994. doi: 10.1007/978-1-4471-2097-1_111.

Ziyao Li. Kolmogorov–Arnold networks are radial basis function networks. arXiv:2405.06721, 2024.

Ziming Liu, Yixuan Wang, Sachin Vaidya, Fabian Ruehle, James Halverson, Marin Soljačić, Thomas Y. Hou, and Max Tegmark. KAN: Kolmogorov–Arnold networks. In *International Conference on Learning Representations (ICLR)*, 2025. URL https://openreview.net/forum?id=Ozo7qJ5vZi.

Giampiero Marra and Simon N. Wood. Practical variable selection for generalized additive models. *Computational Statistics & Data Analysis*, 55(7):2372–2387, 2011. doi: 10.1016/j.csda.2011.02.004.

Raghunathan Ramakrishnan, Pavlo O. Dral, Matthias Rupp, and O. Anatole von Lilienfeld. Big data meets quantum chemistry approximations: The Δ-machine learning approach. *Journal of Chemical Theory and Computation*, 11(5):2087–2096, 2015. doi: 10.1021/acs.jctc.5b00099.

Julien Siems, Konstantin Ditschuneit, Winfried Ripken, Alma Lindborg, Maximilian Schambach, Johannes Otterbach, and Martin Genzel. Curve your enthusiasm: Concurvity regularization in differentiable generalized additive models. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*, 2023.

Daniel Soudry, Elad Hoffer, Mor Shpigel Nacson, Suriya Gunasekar, and Nathan Srebro. The implicit bias of gradient descent on separable data. *Journal of Machine Learning Research*, 19(70):1–57, 2018. URL https://jmlr.org/papers/v19/18-188.html.

Héctor J. Sussmann. Uniqueness of the weights for minimal feedforward nets with a given input-output map. *Neural Networks*, 5(4):589–593, 1992. doi: 10.1016/S0893-6080(05)80037-1.

Simon N. Wood. *Generalized Additive Models: An Introduction with R*. Chapman and Hall/CRC, 2nd edition, 2017. doi: 10.1201/9781315370279.

Yuan Yao, Lorenzo Rosasco, and Andrea Caponnetto. On early stopping in gradient descent learning. *Constructive Approximation*, 26(2):289–315, 2007. doi: 10.1007/s00365-006-0663-2.

Liangwewi Nathan Zheng, Wei Emma Zhang, Lin Yue, Miao Xu, Olaf Maennel, and Weitong Chen. Free-knots Kolmogorov–Arnold network: On the analysis of spline knots and advancing stability. arXiv:2501.09283, 2025.

A All pre-registered predictions

Generated by `F-01` from the committed verdict sections; the cell text is the committed wording. Each study opens with its pre-registration's title and its committed question (or theory). Commits are given as pre-registration → results (in the anonymous version, as neutral labels S-NN). Thresholds are part of each prediction's wording.

A.1 Synthetic controlled study

Pre-registration: "Pre-registration: synthetic teacher-student test of the runaway mechanism (paper2/synth)". **Teacher and students: KAN [4, 3, 1].** Two stacked `kan.KANLayer`s with tanh between them, as in `kan.KAN`, but **without** the salt embedding that `kan.KAN` always adds. Same edge type: SiLU base plus 8 Gaussian RBFs on fixed centres linspace(−1, 1, 8), width 2/7.

ID	Prediction (threshold)	Result	Verdict	Commits
X1	runaway (≥ 10) in correlated P0 and P1; ≤ 3 in the three independent cells and correlated P2	Correlated P0 ×65.7 and P1 ×32.1, and P2 ×1.00, as predicted. **But independent P0 ×12.5 and P1 ×7.9 exceed 3** (independent P2 ×1.09)	**failed**	S-35→S-36
X2	certified $\leq 2/8$ in correlated P0 and P1; $\geq 6/8$ elsewhere	correlated P0 0/8, P1 0/8; independent 6/8, 6/8, 8/8; correlated P2 8/8	**held**	S-35→S-36
X3	cancellation in correlated P0 and P1 each $\geq 2 \times$ correlated P2	P0 3.29 and P1 2.32, against $2 \times 1.99 = 3.98$	**failed**	S-35→S-36
X4	off-manifold ratio ≥ 3 in correlated P0 and P1, ≤ 1.5 in P2	P0 1.22, P1 **0.77**; P2 0.78 (the P2 part held)	**failed**	S-35→S-36
X5	Spearman(cancellation, off-manifold RMSE) ≥ 0.6 over the 24 correlated models	$\rho = +0.177$, p = 0.407	**failed**	S-35→S-36
X6	unit match ≥ 0.85 in the independent cells and correlated P2; < 0.7 in correlated P0 and P1	The first part held: 1.000 in all four no-runaway-predicted cells. **The second part failed: correlated P0 0.774 and P1 0.883 are not < 0.7**	**failed**	S-35→S-36

A.2 Real-data case study

From the case-study close-out: KAN-δ predicts Li-electrolyte conductivity for the CALiSol-23 DEC holdout. Its prediction is log k = a measured pure-PC anchor + x_co · h(solvent descriptors, 1/T, molality, salt). Its first layer is a [8 continuous inputs + 2 salt-embedding, 6, 1] FastKAN, with SiLU + 8 Gaussian RBFs per edge and a smoothness penalty (`F-02`, `F-03`).

Rows are as committed in the case-study summary (one row may cover several predictions, e.g. F2–F5). The rank intervention's verdicts are given per arm in the result column.

Study	Prediction (threshold)	Result	Verdict	Commits
Gauge study	**Control 1**: teacher edges recovered, median dense corr ≥ 0.8	0.301 (sanity: h corr 0.913 < 0.95)	**failed**	S-02→S-03
Gauge study	P1: gauge fixing removes $\geq 50\%$ of edge variance	not run	inconclusive (gated by Control 1)	S-02→S-03
Gauge study	P2: sparse $\geq 2\times$ dense, DEC region highest	not run	inconclusive (gated)	S-02→S-03
Gauge study	P3: (e) < (d) < (c) < (a)	not run	inconclusive (gated)	S-02→S-03
Gauge study	P4: every seed has a unit with $\lvert\text{corr}\rvert < 0.8$	not run	inconclusive (gated)	S-02→S-03
Gauge study	P5: KAN-full edge variance > KAN-δ	not run	inconclusive (gated)	S-02→S-03

Study	Prediction (threshold)	Result	Verdict	Commits
Gauge study	P6: seed spread flags errors	Spearman +0.169 (A) / +0.213 (B); fold B RMSE at 80% coverage 0.153 vs random 0.137	**failed**	S-02→S-04
Control 2	C1: (I,C) edges ≥ 0.9 and units ≥ 0.9	prediction corr 0.947 < 0.99	inconclusive (gate failed)	S-05→S-06
Control 2	C2: (R,C) invT, molal ≥ 0.9; co-solvent < 0.7	prediction corr 0.969 < 0.99	inconclusive (gate failed)	S-05→S-06
Control 2	C3: 107 epochs gives prediction corr < 0.95	0.858 (I,S), 0.919 (R,S)	**held**	S-05→S-06
Control 3	C1': (I) passes the gate, edges ≥ 0.9, units ≥ 0.9	(I) RMSE/std 0.128 > 0.05; unit match 0.648	**failed**	S-07→S-08
Control 3	C2': (R) edges $\leq$ (I) edges $- 0.10$	(R) edges 0.852, units 0.659	inconclusive ((I) failed the gate)	S-07→S-08
Control 3	C4: every continuous input ≥ 0.9 in (I)	–	inconclusive ((I) failed the gate)	S-07→S-08
Fiber	F1: $\geq 50\%$ of seed pairs on-fiber	0/190 (closest d = 0.146)	**failed**	S-09→S-10
Fiber	F2–F5	not run	inconclusive (no on-fiber pairs)	S-09→S-10
Fiber 2	G1: $\geq 4/6$ seeds and $\geq 10/20$ students certify convergence	0/6, 0/20 (all at the 20,000-iteration cap)	**failed**	S-11→S-12
Fiber 2	G2–G5	not run	inconclusive (no converged model)	S-11→S-12
Support	H1: dim(N) ≥ 5	dim(N) = 36 of 72	**held**	S-13→S-14
Support	H2: $\geq 50\%$ of N is REDISTRIBUTION within the correlated group	18/36	**held, exactly at threshold**	S-13→S-14
Support	H3: set A fraction in N ≥ 0.5 and $\geq 3\times$ random	median 0.942 (random 0.5)	inconclusive (untestable: $3 \times 0.5 > 1$)	S-13→S-14
Support	H4: set B (capped), same criterion	median 0.949	inconclusive (untestable)	S-13→S-14
Support	H5: flattest Hessian directions overlap N ≥ 0.5	median 0.885; median layer-1 block share only 0.08	**held**, with caveat (see `support_F-04`)	S-13→S-14
Support	H6: fraction in N rises from control-3 to capped state, $\geq$ 15/20	20/20 (0.713 $\rightarrow$ 0.949)	**held**	S-13→S-14
Rule	I1–I5 (gate)	0 columns removed; gate stopped	inconclusive	S-16→S-17
Rank	I1: median max|grad| $\leq 0.1 \times$ baseline	arm G: **failed** (4.00×); arm F: **failed** (29.1×)	see result	S-18→S-19
Rank	I2: median d_AB $\leq 0.5 \times$ baseline	arm G: **failed** (0.945×); arm F: **failed** (0.808×)	see result	S-18→S-19
Rank	I3 (sanity): median loss $\leq 1.10 \times$ baseline	arm G: passed (1.004×); arm F: passed, narrowly (1.099×)	see result	S-18→S-19
Rank	I4: fold A RMSE ≤ 0.150 (KAN-δ 0.1420)	arm G: **held**, narrowly (0.1499); arm F: **failed** (0.1587)	see result	S-18→S-19
Rank	I5: edge corr on 1000/T and molality $\geq$ baseline + 0.1 (baseline 0.102)	arm G: **held** (0.927); arm F: **held** (0.338)	see result	S-18→S-19
Rank	I6: arm F's improvement $\geq$ arm G's in both I1 and I2	arm G: **failed** (max|grad| 29.1× vs 4.00×; d_AB does favour F, 0.808 vs 0.945)	see result	S-18→S-19
Trust	T1 / T2: trust-krylov certifies $\geq$ 4/6 seeds / $\geq$ 2/3 students	0/6, 0/3. All 9 stopped on scipy's "bad approximation" after 1–62 iterations; the Hessian was accurate to $\leq 2 \times 10^{-7}$	**failed**	S-20→S-21
Trust	T3–T5	nothing converged	inconclusive	S-20→S-21
Stage 2	U1: trust-exact (exact Hessian, 500 iterations) certifies $\geq 4/6$ seeds	0/6. All 9 hit the iteration cap	**failed**	S-22→S-26

Study	Prediction (threshold)	Result	Verdict	Commits
Stage 2	U2: with L2 $\lambda = 10^{-5}$, $\geq 4/6$ seeds and $\geq 2/3$ students	1/6 seeds, 0/3 students	**failed**	S-22→S-26
Stage 2	U3: more converge at 10^{-5} than at 10^{-6}	1 vs 0	**held**	S-22→S-26
Stage 2	U4: converged L2 models' loss without L2 $\leq 1.05 \times$ capped	1.2545	**failed**	S-22→S-26
Stage 2	U5: converged seeds' disagreement $\leq 0.5 \times$ capped	only 1 converged	inconclusive	S-22→S-26
Close	V1: at some μ, $\geq 4/6$ seeds and $\geq 2/3$ students converge (stage2 criterion)	0 certified at every μ	**failed**	S-29→S-30
Close	V2: converged seeds' data term $\leq 1.05 \times$ stage2 A	nothing certified at the selected $\mu = 10^{-6}$	inconclusive	S-29→S-30
Close	V3: $\|\theta\|^2 \leq 10 \times$ start	same	inconclusive	S-29→S-30
Close	V4: seed disagreement halved	same	inconclusive	S-29→S-30
Close	V5: students' unit match ≥ 0.8	same	inconclusive	S-29→S-30
Close	V6: fold A RMSE ≤ 0.150	1.0262, bias +0.9028	**failed**	S-29→S-30
Certify	W1: at $\mu = 10^{-4}$, $\geq 5/6$ seeds certified (second-order criterion, F-a)	6/6; all 17 candidates certified	**held**	S-31→S-32
Certify	W2: certified seeds' median d_AB $\leq 0.5 \times$ capped	0.0771 vs capped 0.1041 (-26%)	**failed**	S-31→S-32
Certify	W3: certified students' unit match ≥ 0.8 (vs the original teacher)	0.624	**failed**	S-31→S-32
Certify	W4: $\mu = 10^{-4}$ fold A RMSE ≤ 0.150	0.3345	**failed**	S-31→S-32
Certify	W5: $\mu = 10^{-4}$ fold A RMSE ≤ 0.5131	0.3345	**held**	S-31→S-32
Certify	W6: Spearman(cross-edge ratio, held-out RMSE) ≥ 0.6	$\rho = +0.345$, p = 0.328 (weak test)	**failed**	S-31→S-32

Committed tally of the case study:

Study	Held	Failed	Inconclusive
First table (Control 1, P1–P6, C1–C3, C1'/C2'/C4, F1–F5, G1–G5, H1–H6)	5	5	19
Rule (I1–I5)	0	0	5
Rank (I1, I2, I4, I5 per arm, I6)	3	6	0
Trust (T1–T5)	0	2	3
Stage 2 (U1–U5)	1	3	1
Close (V1–V6)	0	2	4
Certify (W1–W6)	2	4	0
Total	**11**	**22**	**32**

A.3 Shared-curve KAN

Pre-registration: "SC-KAN-δ (shared-curve KAN-δ): pre-registration". **Inputs:** the 8 continuous F-03 features (flag 11), scaled exactly as in F-03. Salt enters through a learned 2-number embedding e(salt), passed through tanh.

Run after the case-study close-out; not part of the 65 predictions above.

ID	Prediction (threshold)	Result	Verdict	Commits
S1	SC LOCO RMSE $\leq 1.05 \times$ baseline	unweighted mean SC 0.2923 vs base 0.2554 (ratio 1.145; threshold 0.2681). Also reported: row-weighted 0.2478 / 0.2297 (1.079); folds ≥ 20 rows 0.3138 / 0.2093 (1.499)	failed	S-38→S-40
S2	SC fold A RMSE ≤ 0.160	SC 0.1757 (base 0.1420)	failed	S-38→S-40

ID	Prediction (threshold)	Result	Verdict	Commits
S3	SC median d_AB $\leq 0.5 \times$ baseline	SC 0.1865 vs base 0.1120 (threshold 0.0560): SC seeds agree *less* than baseline seeds	failed	S-38→S-40
S4	curve corr ≥ 0.90 for ≥ 6 of 8 inputs	3 of 8 (M_co 0.959, mix_lneta 0.970, invT 0.986); lowest x_co 0.596; others 0.871–0.891	failed	S-38→S-40
S5	$\geq 8/10$ SC fold A models certified	SC 10/10 (max|grad| $\leq 7 \times 10^{-10}$, 0 near-null eigenvalues); baseline 0/10 (all hit the 500-iteration cap)	held	S-38→S-40
S6	median β-norm polished/early ≤ 3 (calibrated)	median 1.337 (range 0.652–2.356); $\|\theta\|^2$ $\times$4.59 (base $\times$67.9, range 30.6–281)	held	S-38→S-40
S7	SC polished/early fold A RMSE $\leq$ 1.5	SC 1.920 (0.1757 $\rightarrow$ 0.3373); baseline 5.736 (0.1420 $\rightarrow$ 0.8147)	failed	S-38→S-40
S8	d log k/d(1000/T) < 0 on $\geq 95\%$ rows in $\geq 8/10$ seeds	10/10 (median share 1.000; 1424 rows, 10 excluded per D4)	held	S-38→S-40
S9	Spearman(1000/T, |slope|) > 0 in $\geq 8/10$ seeds	10/10 (median +0.661, range +0.406 to +0.885)	held	S-38→S-40
S10	d log k/d(ln η_mix) < 0 on $\geq 90\%$ rows in $\geq 8/10$ seeds	7/10 (median share 0.967; s0 0.892, s2 0.892, s8 0.826)	failed	S-38→S-40

A.4 Two-anchor excess model

Pre-registration: "EX-KAN (two-anchor excess model): pre-registration". log k = anchor_PC + x_co · s(eps_co, ln_eta_co, M_co, 1000/T) + x_co · (1 − x_co) · g(eps_co, ln_eta_co, M_co, 1000/T, molality, salt) So δ = log k − anchor_PC = x_co · [s + (1 − x_co) · g].

Run after the case-study close-out; not part of the 65 predictions above.

ID	Prediction (threshold)	Result	Verdict	Commits
E1	EX LOCO $\leq 0.95 \times$ KAN-δ	unweighted mean EX 0.2834 vs threshold 0.2426 (base 0.2554; ratio 1.110). Row-weighted 0.2363 / 0.2297 (1.029); folds ≥ 20 rows 0.2662 / 0.2093 (1.272). EX is better on EC (0.237 vs 0.256), Sulfolane, 3-Glyme, 2-Glyme and Toluene, and worse on 2-MeTHF, THF, Methylene chloride, DME and EA	failed	S-41→S-43
E2	EX fold A ≤ 0.142	EX 0.1502 (base this run 0.1420)	failed	S-41→S-43
E3	s agrees across seeds (separable co-solvents, Amendment 1)	median pairwise corr +0.987 over EA, EC, Methylene chloride: a weak test with 3 values per seed, and Methylene chloride is marginal. Seed SD of s: EA 0.063 (0.39 $\times$ the spread across co-solvents, 0.159), EC 0.091 (0.57 $\times$), Methylene chloride 0.067 (0.42 $\times$). Secondary (≤ 2 rule, 7 co-solvents): +0.869	held	S-41→S-43
E4	mean s(low ε) $<$ mean s(EC, Sulfolane) in $\geq 8/10$ seeds	0/10: the ordering is reversed in every seed. Per high-ε co-solvent: vs EC alone 0/10 (EC median s +0.104), vs Sulfolane alone 0/10 (median −0.124). The low-ε mean is +0.19 to +0.49 per seed. Sulfolane does not drive the failure on its own: EC alone fails in every seed too	failed	S-41→S-43
E5	median seed-SD flagged $\geq 2 \times$ separable (Amendment 1)	0.091 vs 0.067, ratio 1.37. Secondary (≤ 2 rule): 0.222 vs 0.069, ratio 3.22, no verdict; the secondary is driven by 3-Glyme (SD 0.478, M_co extrapolated) and 2-Glyme (0.222)	failed	S-41→S-43

Revision studies. The studies below were run for the revision, each with its own pre-registration.

A.5 Freeze test: is composition necessary? (and weight decay)

Pre-registration: "Revision step 2: is composition necessary? (plus weight decay)". With layer 1 fixed, the output is linear in the layer-2 parameters: a fixed basis times coefficients. The data term plus a quadratic penalty is then a convex quadratic. It can have flat directions but no tilted ones, so it cannot run away. If the runaway needs composition, freezing layer 1 should remove it.

ID	Prediction (threshold)	Result	Verdict	Commits
Z1	Arm A certified: $\geq 5/6$ real; $\geq 7/8$ in each synthetic cell	A-R 6/6; A-S correlated P0 8/8; A-S correlated P1 8/8	HELD	S-48→S-50
Z2	Arm A median layer-2 growth ≤ 3 in R and in each synthetic cell	A-R ×582.8 (×206–×11,689); P0 ×4.07 (×1.36–×245); P1 ×22.9 (×1.51–×801)	FAILED	S-48→S-50
Z3	Arm A smallest Hessian eigenvalue $\geq -10^{-8} \times$ largest at every logged iterate	worst ratio $+1.99 \times 10^{-10}$ (all values positive)	HELD	S-48→S-50
Z4	Arm C: 8/8 certified, median growth ≤ 3	8/8; ×2.52 (×1.68–×2.85)	HELD (computed before the stop)	S-48→S-50
Z5	Arm D branch	8/8 certified; median growth ×0.415 (×0.263–×0.540). Note: this rule uses the same ≤ 3 growth threshold as Z2 and shares its weakness (see the interpretation note); here the norm shrinks, so the threshold is not binding.	"an L2 term of this size at polish time bounds the runaway"	S-48→S-51
Z6 E(i)	Arm E, synthetic correlated P0	median $\|\theta\|^2$ ratio wd 0 / wd 10^{-4} ×1.0011 (×0.987–×1.007); in-distribution RMSE 0.0091 vs 0.0091; off-manifold 0.0667 vs 0.0665; determinism check max |diff| 0	"early stopping alone kept them bounded; weight decay was incidental"	S-48→S-51
Z6 E(ii)	Arm E, real fold A	median ratio ×1.0001 (×1.0001–×1.0001); fold A ensemble RMSE 0.1420 vs 0.1420; determinism check max |diff| 0	"early stopping alone kept them bounded; weight decay was incidental"	S-48→S-51

A.6 Infimum test: a certified regularisation path

Pre-registration: "Revision step 3: infimum test (a certified L2 regularisation path)". Does the unregularised loss of the full model have a finite minimum that is merely distant? Or does its infimum lie at infinity (a receding minimum)?

Predictions I1–I4 apply per phase (Phase 1 synthetic, Phase 2 real data); the real-data control arm was added after Phase 2 with its outcomes fixed before it ran.

ID	Prediction (threshold)	Result	Verdict	Commits
I1	Phase 1: runaway: $\|\theta^*\|^2$ increases as γ decreases	16/16	HELD	S-53→S-54
I2	Phase 1: runaway: $\alpha > 0.2$, bootstrap 95% interval excluding 0	16/16	HELD	S-53→S-54
I3	Phase 1: control: ratio between the two smallest certified $\gamma \leq 1.5$	8/8	HELD	S-53→S-54
I4	Phase 1: runaway: L^* decreases as γ decreases	16/16	HELD	S-53→S-54
I1	Phase 2: runaway: $\|\theta^*\|^2$ increases as γ decreases	6/6	HELD	S-53→S-55

ID	Prediction (threshold)	Result	Verdict	Commits
I2	Phase 2: runaway: $\alpha > 0.2$, bootstrap 95% interval excluding 0	6/6	HELD	S-53→S-55
I3	Phase 2: discrimination check	—	Phase 1's I3 (HELD, 8/8), as pre-registered; Phase 2 has no control models	S-53→S-55
I4	Phase 2: runaway: L^* decreases as γ decreases	6/6	HELD	S-53→S-55
Control	Post hoc, outcomes fixed before running. Plateau: the $\|\theta^*\|^2$ ratio between the two smallest certified (incl. certified-degenerate) γ values is ≤ 1.5 for both seeds. The real-data contrast is then complete.	The $\|\theta^*\|^2$ ratio between the two smallest certified γ is 1.092 (ind0) and 1.101 (ind1), both ≤ 1.5, so the real-data contrast is complete. The divergence condition is not met: α is +0.105 [+0.093, +0.118] and +0.124 [+0.109, +0.139], both below 0.2.	plateau	S-55→S-56

A.7 Correlation dose-response

Pre-registration: "Revision step 4: dose-response — how much correlation does the blind spot need?". At what input correlation does the runaway turn on, and does the null-space penalty still close it at the strongest dose?

Verdict column: P0 / P1. The thresholds are synthetic ([4, 3, 1], Gaussian inputs), not real-data thresholds.

ID	Prediction (threshold)	Result	Verdict	Commits
R1	The uncertified fraction (measure 1) is non-decreasing across ρ, for P0 and for P1 separately, allowing one inversion.	P0: HELD: 0.125, 0.625, 0.500, 1, 1, 1 (1 inversion); P1: HELD: 0.125, 0.500, 0.500, 0.625, 1, 1 (0 inversions)	held / held	S-57→S-61
R2	Control, P2 at $\rho = 0.99$: at least 7 of 8 certified (including certified-degenerate), and median D ≤ 1.5.	P2 at $\rho = 0.99$: 8/8 certified, median D 1.127 (range of growth ×0.88–×1.11).	held	S-57→S-61
R3	The median affine-RBF share of the polished change is non-decreasing in ρ, for P0 and for P1 (one inversion allowed).	P0: FAILED: 0.715, 0.594, 0.612, 0.476, 0.394, 0.585 (3 inversions); P1: FAILED: 0.944, 0.982, 0.748, 0.934, 0.914, 0.952 (2 inversions)	failed / failed	S-57→S-61
R4	Median D increases with ρ, for P0 and for P1: Spearman correlation ≥ 0.8 between ρ and the cell median D over the 6 levels. (Spearman ≥ 0.8 is the operative criterion; it tolerates up to about two inversions.)	P0: FAILED: D 8.01, 15.05, 6.43, 9.52, 16.13, 16.72; Spearman +0.714; P1: FAILED: D 4.54, 4.88, 2.91, 3.00, 4.52, 14.54; Spearman +0.143	failed / failed	S-57→S-61
	R4, certified-only variant	P0: FAILED (Spearman +0.714; D at 0.99 is 19.55, n = 4); P1: FAILED (Spearman +0.143)	failed / failed	S-57→S-61

A.8 B-spline replication

Pre-registration: "Revision step 5: B-spline replication of the synthetic 2 × 3 study". Are the runaway and its fix a property of KAN-style edges with second-difference penalties, or an artefact of the Gaussian-RBF basis?

Step 1 verdicts at the frozen $\mu = 10^{-6}$, then the post-hoc follow-up checks whose predictions and readings were fixed before computing (check 3 is descriptive; Appendix B).

ID	Prediction (threshold)	Result	Verdict	Commits
B1	Correlated P0 and correlated P1: at most 2 of 8 certified, each.	correlated P0: 0/8 (8 iteration cap)	HELD	S-64→S-79
		correlated P1: 0/8 (8 iteration cap)	HELD	S-64→S-79
B2	P2, both input conditions: at least 7 of 8 certified AND median D $\leq$ 1.5, each.	independent P2: 8/8, D 1.222 (certified-only 1.222)	HELD	S-64→S-79
		correlated P2: 8/8, D **2.225** (certified-only 2.225)	FAILED (the D leg; certification held)	S-64→S-79
B3	Independent P0 and P1: at least 6 of 8 certified, each (the Gaussian reference level).	independent P0: 5/8 (3 iteration cap)	FAILED	S-64→S-79
		independent P1: 4/8 (3 iteration cap, 1 solver error via the watchdog)	FAILED	S-64→S-79
Check 1	Post hoc, predictions fixed in advance; P2 path extended to $\gamma = 10^{-9}$: median D' $\leq$ 1.5 in both conditions;; median norm ratio $\|\theta^*(10^{-9})\|$ / $\|\theta_\text{polish}\|$ within [0.9, 1.1] in both conditions.	independent: D' 1.003, norm ratio 1.000; correlated: D' 1.035, norm ratio 0.911	both hold	S-80→S-83
Check 2	Post hoc, prediction fixed in advance; P2 at $\mu = 10^{-5}$: $\geq$ 7/8 certified and median D $\leq$ 1.5 in both conditions.	independent: 8/8 certified, median D 1.167 (1.167); correlated: 8/8 certified, median D 1.309 (1.309)	held / held	S-80→S-83

B Post-hoc analyses

Every analysis below except one was defined after seeing results; the exception, the applicability analysis, is descriptive, with its plan committed before computing. Each row's label states how the analysis was defined, as recorded in its own files; none changes a pre-registered verdict, and all are reported as exploratory. The description is the subject line of the commit that added the analysis (plan and results commits where both exist).

Analysis	Commit subject (as committed)	Commit
Post hoc. Anatomy of the real-data runaway `anatomy_F-05, F-06, anatomy_F-07`	Plan for descriptive runaway anatomy (before F-06 exists)	S-27
	Anatomy (descriptive): Q1, Q2 not as expected, Q3 as expected; RBF growth is in the penalty null space	S-28
Post hoc. Gauge check and regularised teacher (exploratory) `F-08, gaugecheck_F-09, gaugecheck_F-07`	gaugecheck (POST-HOC, exploratory): penalty fixes gauge position by slope balance; students match the regularised teacher	S-33
Post hoc. Anatomy of the synthetic runaway `paper2/anatomy_F-10, paper2/anatomy_synth_F-07`	paper2 anatomy_synth (POST-HOC, descriptive): synthetic runaway lives mainly in layer-2 output-edge cancellation, unsaturated; no verdicts change	S-37
Post hoc. Loss decomposition along the runaway `paper2/F-11, paper2/lossdecomp_F-07`	paper2 lossdecomp (POST-HOC, descriptive): synthetic runaway loss decrease shared, D2 penalty median 0.61 vs data 0.37; 12/20 runaway models overfit; real-data reference recomputed	S-44
Post hoc. Descriptive checks (hidden-unit correlation, tilt affinity, references) `paper2/checks/F-05, paper2/checks/F-12, paper2/checks/checks_F-07`	paper2/checks: plan for descriptive checks (F4 hidden-unit correlation, F2 tilt affinity, A1 commit and data citation) before F-12 exists	S-45

Analysis	Commit subject (as committed)	Commit
	paper2/checks (POST-HOC, descriptive): hidden-unit correlation marks the runaway (independent cells 0.70 vs 0.05) but is not sufficient; RBF tilt affine to 2% centrally, 13% at the ends; A1 commit S-01; CALiSol-23 citation not in competition material	S-46
Post hoc. Degeneracy follow-up: restart, scaling and flat-direction probes of the freeze test's endpoints `paper2/revision/F-13, paper2/revision/degeneracy_F-07`	Post-hoc degeneracy follow-up: crash record 2 (trust-exact subproblem hang in the B-R ind0 restart probe; timer did not fire) and implementation-only fix (restart probes in a killable child process with checkpoints)	S-52
Post hoc. Infimum test, Phase 0: re-classification of the freeze test's degenerate endpoints `paper2/revision/infimum_F-04`	Infimum test Phase 0 (POST-HOC, no class changed) and Phase 1 (I1-I4 held 16/16, 16/16, 8/8, 16/16: receding minimum over gamma 10^{-2}..1e-7); post-hoc cold-start basin check (no small-norm basin; better larger-norm minima at small gamma); verdicts appended	S-54
Post hoc. Infimum test: cold-start basin check `paper2/revision/infimum_F-04`	Infimum test Phase 0 (POST-HOC, no class changed) and Phase 1 (I1-I4 held 16/16, 16/16, 8/8, 16/16: receding minimum over gamma 10^{-2}..1e-7); post-hoc cold-start basin check (no small-norm basin; better larger-norm minima at small gamma); verdicts appended	S-54
Post hoc, outcomes fixed in advance. Infimum test: real-data control arm (null-space penalty at $\mu = 10^{-4}$) `paper2/revision/infimum_F-04`	Infimum test: Phase 2 verdicts; pre-register POST-HOC real-data control arm (before it runs)	S-55
	Infimum test Phase 2 (I1, I2, I4 held 6/6; receding minimum over gamma 10^{-2}..1e-7; six distinct branches), arms (a)/(c) (descriptive, n = 3), POST-HOC real-data control arm (mu = 10^{-4} fixed: PLATEAU, ratios 1.09/1.10), cold-start summary; final run notes (interruption = session end, cumulative cap held; Phase 1 equivalence hashes)	S-56
Post hoc, rule fixed in advance. Dose-response settling check (short path against the certified polished minimum) `paper2/revision/dose_F-14, paper2/revision/dose_posthoc_F-15`	F-16: POST-HOC descriptive check (short path vs certified polished minimum), rules fixed before computing	S-62
	Dose-response POST-HOC check: 30/32 certified-polish models read FINITE (11 valley, 19 not yet reached), 0 RECEDING, 2 unclassified -> reading A below Branch 5; run notes a-d	S-63
Post-outcome diagnostic, criteria fixed in advance. B-spline null-space diagnostic (constant-shift symmetry) `paper2/revision/bspline_nullspace_F-17, paper2/revision/bspline_nullspace_F-18`	F-19: DESCRIPTIVE constant-shift null-space diagnostic (criteria fixed before results)	S-69
	B-spline null-space diagnostic (DESCRIPTIVE): exact 15-dim constant-shift symmetry confirmed (exactness, containment, RBF contrast, hangs after convergence); criterion (ii) fails at the two correlated end points (16 and 3 extra near-flat, non-exact directions at 10^{-10}..1e-8); stop	S-70
Post hoc, predictions fixed in advance. B-spline follow-up: path extension, $\mu = 10^{-5}$, the non-certifying independent models `paper2/revision/bspline_F-04, paper2/revision/bspline_F-14, paper2/revision/bspline_posthoc_F-15`	B-spline prereg: POST-HOC addendum (P2 path extension, P2 at mu = 10^{-5}, non-certifying independent models), predictions and readings fixed before computing	S-80
	B-spline POST-HOC follow-up: check 1 both predictions hold (D' 1.003 / 1.035; norm ratio 1.000 / 0.911), check 2 holds at mu = 10^{-5} (8/8, D 1.167 / 1.309), check 3 2 runaway-like / 4 mixed / 1 slow; archive 18	S-83

Analysis	Commit subject (as committed)	Commit
Descriptive, plan committed first. Applicability analysis (no predictions) `paper2/revision/appdomain_F-05`, `paper2/revision/F-20`, `paper2/revision/appdomain_F-07`	Applicability-domain plan (DESCRIPTIVE; definitions and figures fixed before computing; facts of record; G1-G4)	S-84
	Applicability domain (DESCRIPTIVE): fold ensembles reproduced (0.1420 / 1.0262 / 0.3345); off-pattern tracks LOCO RMSE better than general distance but n = 7, not significant; within fold A polishing hurt most near the training data, collinear with x_co; fold A on-pattern, DEC 7x further off-pattern; parked study's precondition not met	S-86
Post hoc. Applicability: artefact check (the ensembles' own errors against distance) `paper2/revision/appdomain_F-15`	Applicability domain: close-out; parked predictive study dropped (precondition not met)	S-87
Post hoc. Regularised teacher (weights pinned)	Pin gaugecheck regularised-teacher weights in a seventh archive; append to weights manifest	S-34

C Reproducibility map

Generated from `F-21`. Every figure is rebuilt by its script from committed results; each figure's sidecar JSON lists the data files (path, commit, SHA-256), any archived weights (SHA-256, archive) and every plotted number. Code will be released upon publication.

Figure	Script	Data files (commit)	Archived weights
Figure 1	`make_F-22`	none (computed from committed code: F-02, F-23)	—
Figure 6	`make_F-24`	`paper2/synth_F-15` (S-36)	—
Figure 7	`make_F-25`	`weights_F-26` (S-85); `paper2/anatomy_synth_F-15` (S-37); `paper2/checks/checks_F-15` (S-46)	18 files, archives W-01, W-02, W-03
Figure 8	`make_F-27`	`paper2/revision/appdomain_F-15` (S-87); `weights_F-26` (S-85)	15 files, archives W-04, W-05, W-06
Figure 2	`make_F-28`	`paper2/revision/infimum_F-15` (S-56)	—
Figure 3	`make_F-29`	`paper2/revision/dose_F-15` (S-61)	—
Figure 4	`make_F-30`	`paper2/synth_F-15` (S-36); `paper2/revision/bspline_F-15` (S-78); `paper2/revision/bspline_posthoc_F-15` (S-83)	—
Figure 9	`make_F-31`	`stage2_F-15` (S-26); `close_F-15` (S-30); `certify_F-15` (S-32)	—
Figure 5	`make_F-32`	`arch/sc_F-15` (S-40); `certify_F-15` (S-32); `weights_F-26` (S-85)	40 files, archives W-05
Figure 10	`make_F-33`	`paper2/lossdecomp_F-15` (S-44)	—
Figure 11	`make_F-33`	`gaugecheck_F-15` (S-33)	—
Figure 12	`make_F-33`	`arch/sc_F-15` (S-40); `arch/ex_F-15` (S-43)	—

Tables in the main text (no script; values typed in the draft):

Table	Source	Data files	Commits
Table 1	typed in the draft from	`F-34`, `F-35`	S-36
Table 2	typed in the draft from	`F-36`, `F-37`, `F-38`	S-30, S-32, S-26
Table 17	typed in the draft from	`F-39`, `F-40`, `F-41`	S-40, S-43, S-47

Panels evaluated post hoc from archived weights (from `F-42`):

- **fig3a, real data:** the split of the parameter change into affine and curved RBF parts per layer, for 107 epochs → fiber2 capped → stage2 Part A.

- `stage2_F-15` stores only per-block norms, without an affine/curved split and without the 107 → capped change.
- **fig9b, polished fold A d_AB** for SC-KAN-δ and the baseline (marked *).
- arch/sc committed d_AB only at the early state.
- **fig4a, b, fold A per-row error change** (early-stopped -> mu = 10^{-6} and mu = 10^{-4} ensembles).
- F-43 stores only the per-row Spearman correlations and medians; the per-row values are re-evaluated with the

D Research process

Pre-registration. Every study committed a pre-registration file (predictions, thresholds, interpretation rules) before the script that tested it existed. Verdicts were appended below the committed text; the pre-registered text was never edited, and corrections to reported numbers were made in separate, labelled commits.

Amendments. A change to a design after its pre-registration was committed as a dated amendment section before any affected model was trained, with its reason; the original text was kept. Choices forced by the code or data at run time were logged as numbered discrepancies at the top of each output log.

Files win. Instructions to the agent stated that committed files take precedence over any number or description in a prompt; discrepancies between a brief and the files were flagged rather than resolved silently.

Data and weights. Model weights were never committed: they were archived outside the repository with a manifest of SHA-256 hashes and producing commits, and every later analysis verified the hashes and held its inputs in memory. Post-hoc analyses were labelled as such in their files and changed no verdict.

AI assistance. Code, runs and logs were produced by an AI coding agent working under this protocol. Study designs, briefs and the manuscript were developed with an AI assistant. The author set the research questions, reviewed and approved every brief and amendment, decided each next stage, and takes responsibility for all content. Crashes were recorded in the logs and fixed in separate commits labelled as implementation-only.

Pre-registrations, amendments and crash records (from the commit history):

Commit	Date	Kind	Commit subject
S-02	2026-09-23	pre-registration	Pre-register KAN identifiability study (before any experiment)
S-05	2026-09-23	pre-registration	Pre-register corrected 2x2 control (before writing or running it)
S-07	2026-09-23	pre-registration	Pre-register control 3: reachable teacher (before writing code)
S-09	2026-09-29	pre-registration	Pre-register fiber completeness test (before writing or running F-44)
S-11	2026-09-29	pre-registration	Pre-register fiber2: certified convergence + teacher-known fiber analysis (before writing code)
S-13	2026-09-29	pre-registration	Pre-register support study: near-null space, where models differ, Hessian (before F-45 exists)
S-15	2026-09-29	pre-registration	Close-out: electrolyte case study summary (all pre-registered predictions)
S-16	2026-09-29	pre-registration	Pre-register basis-size rule intervention (before F-46 exists)
S-18	2026-09-29	pre-registration	Pre-register rank-constrained layer-1 intervention (before F-47 exists)
S-20	2026-09-30	pre-registration	Pre-register trust-region Newton polishing, stage 1 (before F-48 exists)
S-22	2026-09-30	pre-registration	Pre-register stage 2: trust-exact check, runaway diagnostic, coercivity (before F-49 exists)
S-23	2026-09-30	crash record / fix	F-49 as first run (crashed with AttributeError in callback)
S-24	2026-09-30	crash record / fix	stage2: fix AttributeError (implementation only, post-prereg)
S-25	2026-10-02	crash record / fix	stage2: hold verified inputs in memory (implementation only, post-prereg)
S-29	2026-10-02	pre-registration	Pre-register targeted null-space penalty intervention (before F-50 exists)

Commit	Date	Kind	Commit subject
S-31	2026-10-03	pre-registration	Pre-register second-order certification and regularised fold A accuracy (before F-51 exists)
S-35	2026-10-03	pre-registration	paper2: pre-register synthetic 2x3 runaway test (before F-10 exists)
S-38	2026-10-03	pre-registration	arch: pre-register SC-KAN-δ (shared-curve KAN-δ) before any script exists; decisions C (centring), mu null-space term, tanh salt embedding, full beta term
S-39	2026-10-03	crash record / fix	arch sc_run: fix log() call (crash record 1, after the first fold A polish); resume from cache
S-41	2026-10-04	pre-registration	arch: pre-register EX-KAN (two-anchor excess model) before any script exists; decision B (x_co-independent g, [7,6,1]), full set for E3-E5, E4 mean rule
S-42	2026-10-04	amendment	arch ex_prereg: Amendment 1 (after Step 0, before training): primary separability rule by fixed-composition groups (EA, EC, Methylene chloride separable); committed rule reported as secondary
S-48	2026-10-16	pre-registration	Revision step 2: pre-register the composition test (frozen layer 1, one-layer control) and the weight-decay arms D/E, before F-52 exists; draft rev 37
S-49	2026-10-17	pre-registration	Revision step 2: decision record - F-52 run stopped by the pre-registered rule (Z2 failed in Arm A); the author decided to resume (option 2)
S-50	2026-10-17	crash record / fix	Revision step 2: crash record 1 (trust-exact subproblem hang in Arm B-R) and implementation-only fix (timer-enforced 2 h cap); part-1 verdicts and post-hoc note in the prereg
S-52	2026-10-18	crash record / fix	Post-hoc degeneracy follow-up: crash record 2 (trust-exact subproblem hang in the B-R ind0 restart probe; timer did not fire) and implementation-only fix (restart probes in a killable child process with checkpoints)
S-53	2026-10-18	pre-registration	Pre-register infimum test: certified L2 regularisation path (before F-53 exists)
S-55	2026-10-19	pre-registration	Infimum test: Phase 2 verdicts; pre-register POST-HOC real-data control arm (before it runs)
S-57	2026-10-20	pre-registration	Pre-register dose-response study (before F-54 exists)
S-58	2026-10-20	crash record / fix	F-54 as it ran the replication gate (attempt 1, stopped by crash record 1)
S-59	2026-10-20	crash record / fix	F-54: crash record 1 (SIGBUS in Accelerate dgemm during a polish Hessian) and implementation-only fix (a committed polish restarts from the early state; stale heartbeat discarded)
S-60	2026-10-20	crash record / fix	dose/infimum: crash record 2 (non-finite trust-exact subproblem in scipy) and implementation-only fix (stop at the last accepted iterate with reason 'solver error')
S-64	2026-10-22	pre-registration	Pre-register B-spline replication (before F-55 exists)
S-65	2026-10-22	crash record / fix	B-spline basis (efficient-kan-style extended uniform knots, G = 5, k = 3) and property tests (a)-(e), implementation only
S-66	2026-10-22	amendment	B-spline prereg: amendment after Step 0 (i) failure; fidelity gate withdrawn as mis-specified (before any further run)
S-68	2026-10-22	crash record / fix	bspline: crash record 1 (Step 0 (ii) teacher polish hung in scipy's trust-exact subproblem) and implementation-only fix (teacher polishes in hard-killed children with checkpoints; ending classification per S-67; mu calibration as its own stage)
S-71	2026-10-23	amendment	B-spline prereg: POST-HOC amendment after the null-space diagnostic (analytic 15-dim constant-shift null set; null-set-aware certification, degeneracy count and probes; certification-based termination; B1-B3 unchanged)
S-72	2026-10-23	amendment	B-spline: null-set-aware certification, complement degeneracy count and probes, certification-based termination (amendment S-71); re-certification and validation stages; tests
S-74	2026-10-23	crash record / fix	B-spline: staleness watchdog (kill a child with no new accepted iterate for 10 min; 'solver error (hang)'), implementation-only crash handling (note S-73); tests
S-75	2026-10-23	crash record / fix	B-spline speed-up, implementation only: layer-1 basis cache for fixed inputs + vectorised exact Hessian (switchable)
S-76	2026-10-23	crash record / fix	F-55: the mu-calibration stage reads the amended Step 0 (ii) gate (S-71 item 4); implementation only
S-77	2026-10-23	amendment	B-spline: frozen mu = 10^{-6} (dated amendment; growth and certification criteria agree); Step 1 estimate 1.2-3 h typical, 16 h worst case (below the 3-day checkpoint)
S-81	2026-10-24	crash record / fix	F-55: model keys may carry an explicit P2 mu ('<cond>_P2m<mu>_s<seed>'), implementation only (POST-HOC addendum S-80, check 2)

Commit	Date	Kind	Commit subject
S-82	2026-10-24	crash record / fix	bspline_posthoc: crash record (thread deadlock in libomp during check 1) and implementation-only fix (one worker process per model)

Changes made after an outcome was known

Generated from the committed files: each entry quotes the committed wording, with the date and commit of the quoted text.

Change	Committed wording	Date	Commit
B-spline replication, calibration step (Step 0)	the fidelity gate, withdrawn	2026-10-22	S-66
B-spline replication, calibration step (Step 0)	the polish-ending classification	2026-10-22	S-67
B-spline replication, calibration step (Step 0)	the analytic null set, null-set-aware certification and certification-based termination	2026-10-23	S-71
B-spline replication, calibration step (Step 0)	the staleness watchdog	2026-10-23	S-73
B-spline replication, calibration step (Step 0)	the frozen μ	2026-10-23	S-77
Dose-response: runtime correlation tolerance	**(a)** The runtime correlation tolerance (0.02) was miscalibrated: it is below the sampling error at n = 2000 (about 0.022 at $\rho = 0$). Pairing and the x_1/x_3 identity passed. Future tolerances are derived from the sampling error of the quantity checked.	2026-10-21	S-63
Dose-response: added stop reason	**(b)** "Solver error (trust-exact subproblem)" is a dated addition (2026-10-20) to the frozen stop-reason list (crash records: S-59, S-60).	2026-10-21	S-63
B-spline replication: added stop reason (hang)	A polish killed after hanging records "solver error (hang)" and the fact that it was ended by the 2 h + 5 min child kill. This is consistent with the frozen failure-reason list (S-60).	2026-10-22	S-67
B-spline replication: added stop reason (watchdog)	**Recorded as** "solver error (hang); ended by the staleness watchdog", with the same classification as a hard kill.	2026-10-23	S-73
Synthetic study: X6, pre-registered reading did not fit	**X6's rule maps failure to "non-identifiability persists beyond the runaway". The data contradict that reading.** X6 failed only because the *runaway* cells matched the original teacher *better* than predicted (0.774 and 0.883 against < 0.7). Every cell without a runaway recovered the reference's internals exactly (unit and edge match 1.000, d_AB about 10^{-4}). The verdict stays **failed**, as recorded.	2026-10-03	S-36
Freeze test: Z2, branch taken and run resumed	By the fixed interpretation, Z2 failing selects the branch "a runaway without composition: the mechanism needs revision; stop and report before anything else". The run was stopped and reported. On 2026-10-17 the author decided, after seeing Arm A's results, to resume Arms D, E and B (option 2).	2026-10-17	S-50
Freeze test: Z2, pre-registered reading did not fit	**Z2's branch premise did not occur.** That premise is "a runaway without composition". Every Arm A model reached a certified, finite minimum with non-negative curvature (Z3). The failure reflects the growth threshold, which did not separate a distant finite minimum from unbounded growth. The pre-registered rules did not anticipate this.	2026-10-17	S-50
Applicability analysis: predictive follow-up dropped	The parked predictive study is DROPPED: its pre-stated precondition (off-pattern distance bearing on the fold A polish damage) is not met. Within fold A, the co-solvent off-pattern distance is constant, and every per-row distance is collinear with x_co, through which the damage scales by the model's x_co * h structure. No successor study.	2026-10-25	S-87

E Supplementary figures

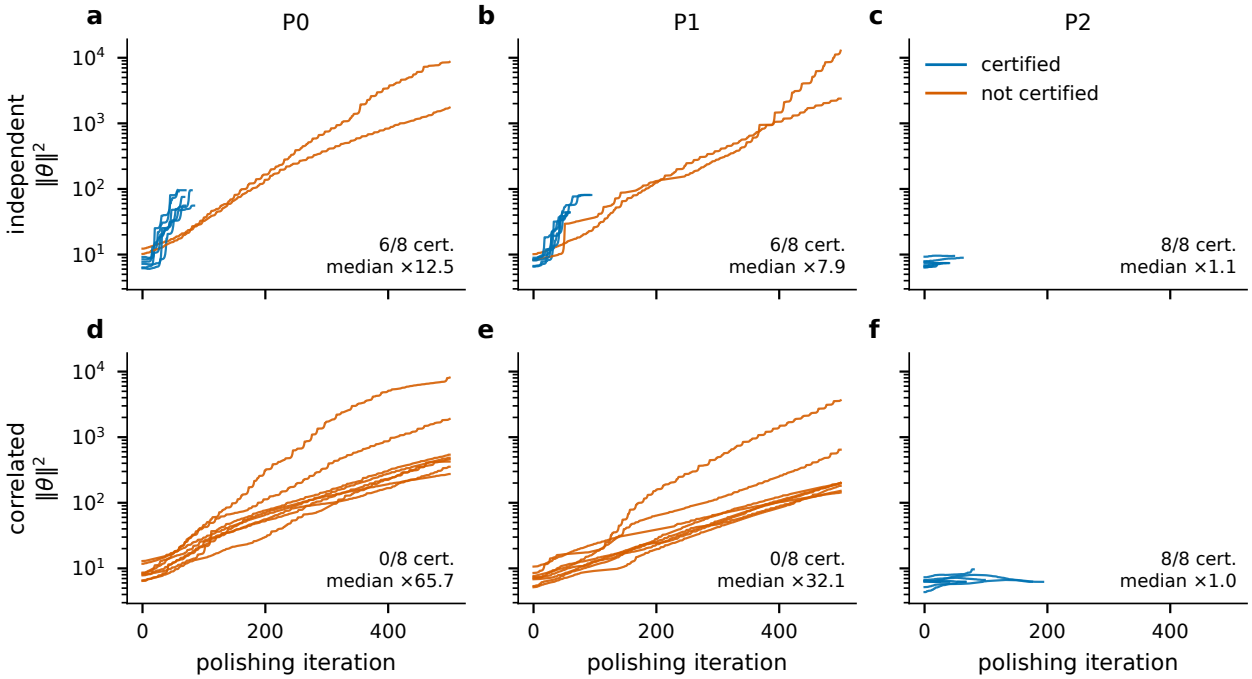


Figure 6: Synthetic 2×3 study: $\|\theta\|^2$ over exact-Hessian trust-region polishing iterations (log scale), one line per seed, coloured certified (blue) or not certified (vermillion). Rows: independent and correlated inputs; columns: P0, P1, P2. Per panel, certified seeds and median growth (polished over early-stopped): independent P0: 6/8 certified, ×12.5; independent P1: 6/8 certified, ×7.9; independent P2: 8/8 certified, ×1.1; correlated P0: 0/8 certified, ×65.7; correlated P1: 0/8 certified, ×32.1; correlated P2: 8/8 certified, ×1.0.

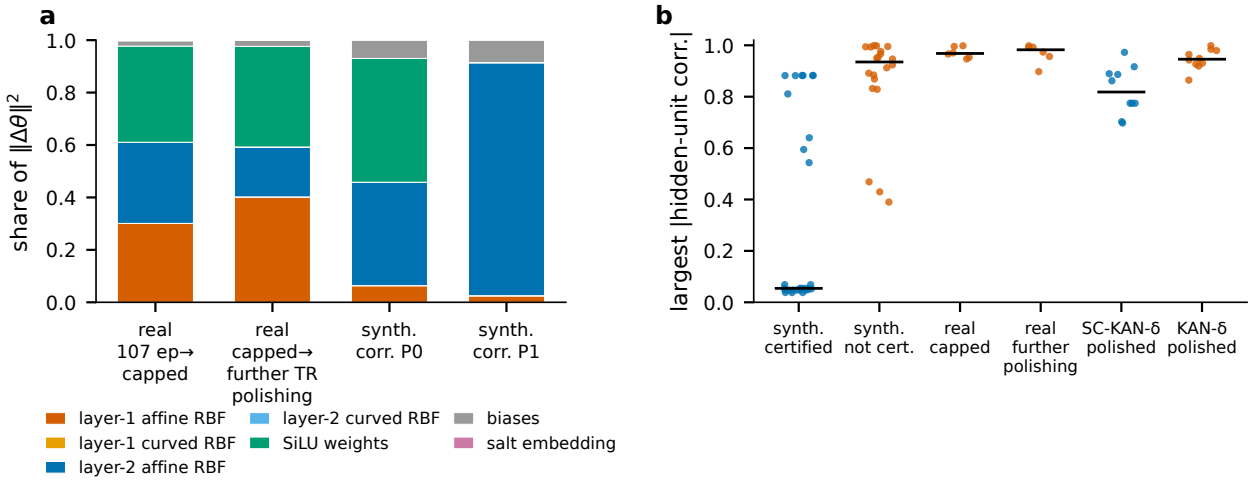


Figure 7: Where the runaway lives. **(a)** Mean share of $\|\Delta\theta\|^2$ by parameter block. Real data, 107 epochs→capped: layer-1 affine RBF 0.30, layer-2 affine RBF 0.31, SiLU weights 0.37; capped→further trust-region polishing: 0.40, 0.19, 0.38; synthetic correlated P0: layer-2 affine RBF 0.39, SiLU weights 0.47; correlated P1: layer-2 affine RBF 0.89. The real-data split was evaluated post hoc from archived weights (Appendix C). **(b)** Largest absolute correlation between hidden-unit activations on the training rows, median: synthetic, certified 0.054 ($n = 28$); synthetic, not certified 0.935 ($n = 20$); real, capped 0.968 ($n = 6$); real, further polishing 0.983 ($n = 6$); SC-KAN-δ polished 0.818 ($n = 10$); KAN-δ polished 0.946 ($n = 10$). Blue: certified; vermillion: not certified or runaway.

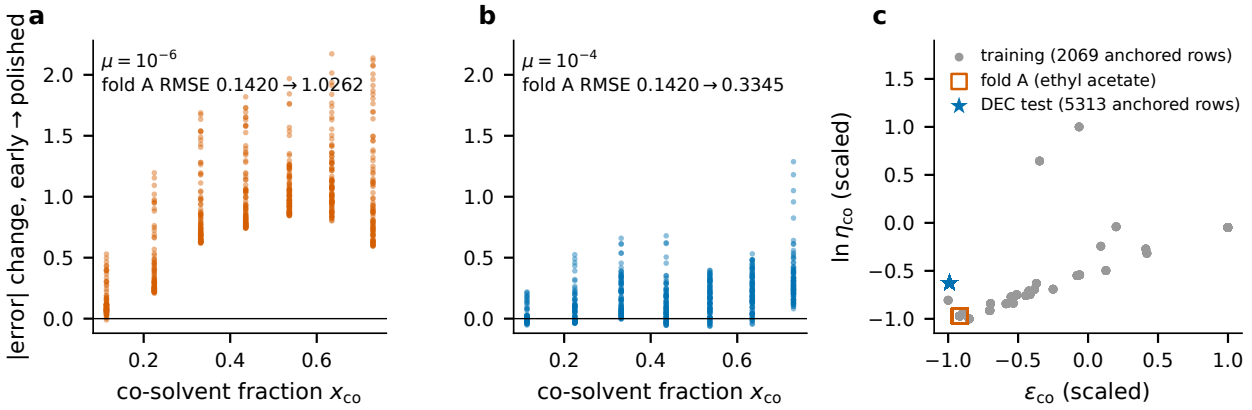


Figure 8: Applicability (descriptive; analysis plan committed before computing). **(a, b)** Fold A (ethyl acetate held out, 635 rows): change in each row's absolute error from the early-stopped 5-seed ensemble to the polished ensemble, against the co-solvent fraction x_{co} (7 distinct values); fold A RMSE 0.1420→1.0262 at $\mu = 10^{-6}$ (a) and 0.1420→0.3345 at $\mu = 10^{-4}$ (b). Per-row values are evaluated post hoc from SHA-verified archived weights, with the logged RMSEs reproduced (Appendix C). **(c)** Co-solvent descriptors ε_{co} and $\ln \eta_{co}$ in the committed min–max scaling: the 2069 anchored training rows (30 distinct points), fold A's co-solvent (off-pattern distance 0.047) and the 5313 anchored DEC test rows (0.334); each held-out set occupies a single point.

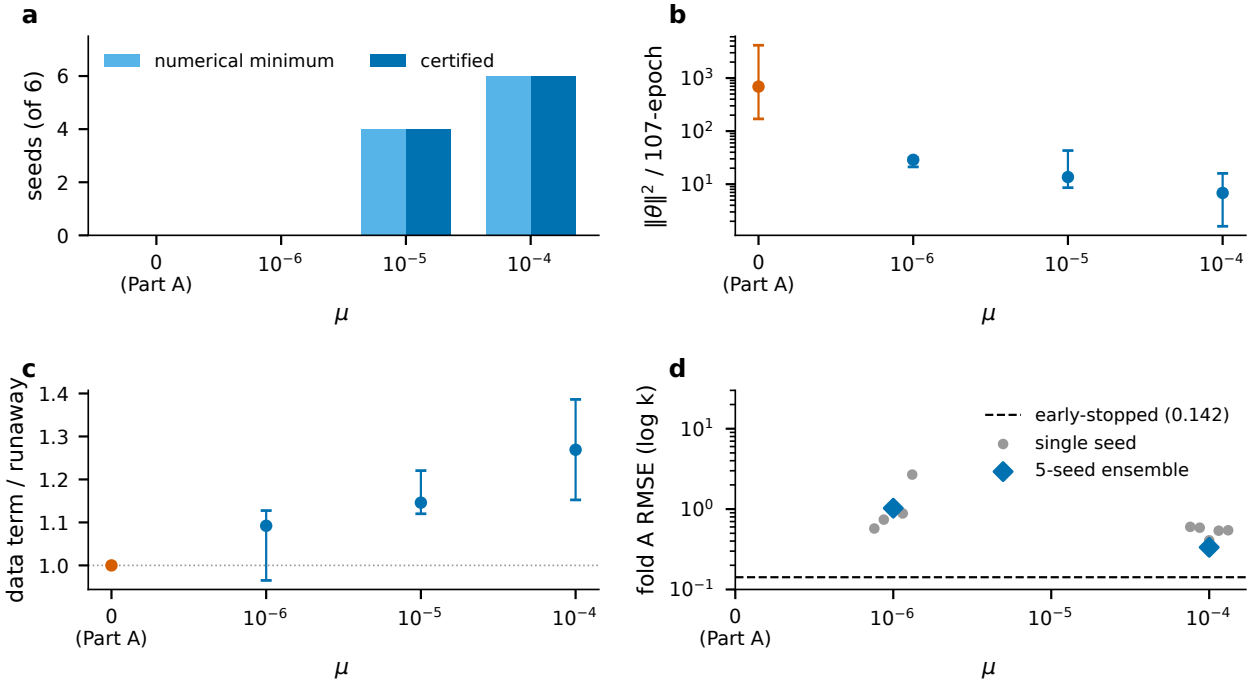


Figure 9: Real-data dose-response in the null-space weight μ (6 KAN-δ seeds; $\mu = 0$: polished without a null-space penalty). **(a)** Seeds at a numerical minimum (all of them certified): $\mu = 0$, 0 of 6; 10^{-6}, 0 of 6; 10^{-5}, 4 of 6; 10^{-4}, 6 of 6. **(b)** $\|\theta\|^2$ relative to the 107-epoch state, median (range): $\mu = 0$, 690 (170–4146); 10^{-6}, 29 (21–31); 10^{-5}, 14 (9–43); 10^{-4}, 7 (2–16). **(c)** Training data term relative to the runaway, median: 10^{-6}, 1.092; 10^{-5}, 1.146; 10^{-4}, 1.269. **(d)** Fold A RMSE: early-stopped 0.142 (dashed line); polished 5-seed ensemble 1.026 at $\mu = 10^{-6}$ and 0.334 at $\mu = 10^{-4}$; grey points are single seeds.

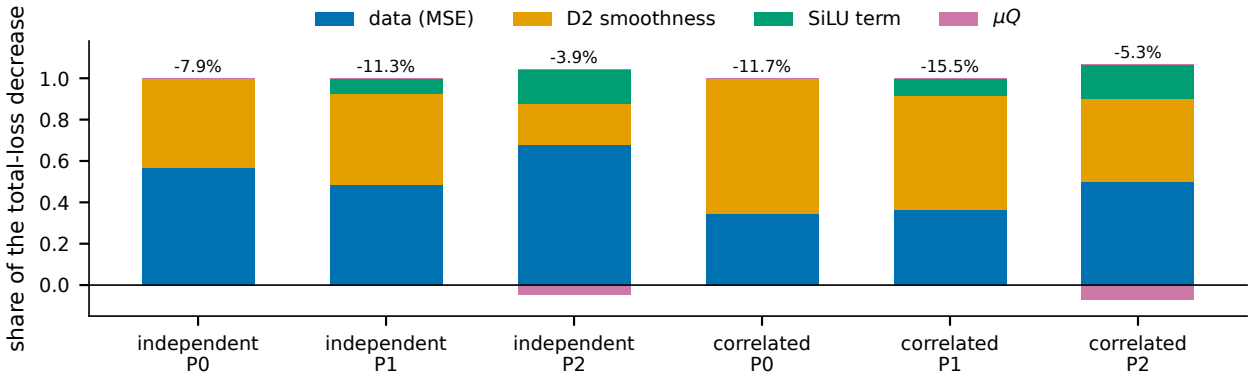


Figure 10: Post hoc loss decomposition (synthetic, early→polished): share of the total-loss decrease from each term, pooled over the 8 seeds of each cell. independent P0: data 0.57, D2 0.43 (median total change -7.9%); independent P1: data 0.49, D2 0.44 (median total change -11.3%); independent P2: data 0.68, D2 0.20 (median total change -3.9%); correlated P0: data 0.35, D2 0.65 (median total change -11.7%); correlated P1: data 0.37, D2 0.55 (median total change -15.5%); correlated P2: data 0.50, D2 0.40 (median total change -5.3%).

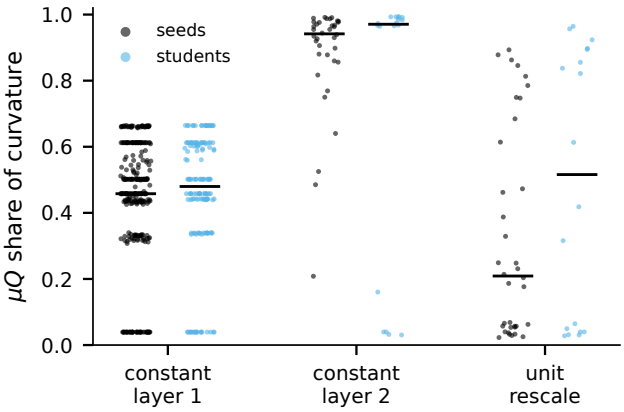


Figure 11: Post hoc gauge check at the $\mu = 10^{-4}$ certified minima: share of the curvature along each least-squares gauge direction contributed by μQ. Medians: constant shift in layer 1, real-data seeds 0.46 ($n = 288$); constant shift in layer 1, students (trained on a frozen teacher) 0.48 ($n = 144$); constant shift in layer 2, real-data seeds 0.94 ($n = 36$); constant shift in layer 2, students 0.97 ($n = 18$); unit rescaling, real-data seeds 0.21 ($n = 36$); unit rescaling, students 0.52 ($n = 18$).

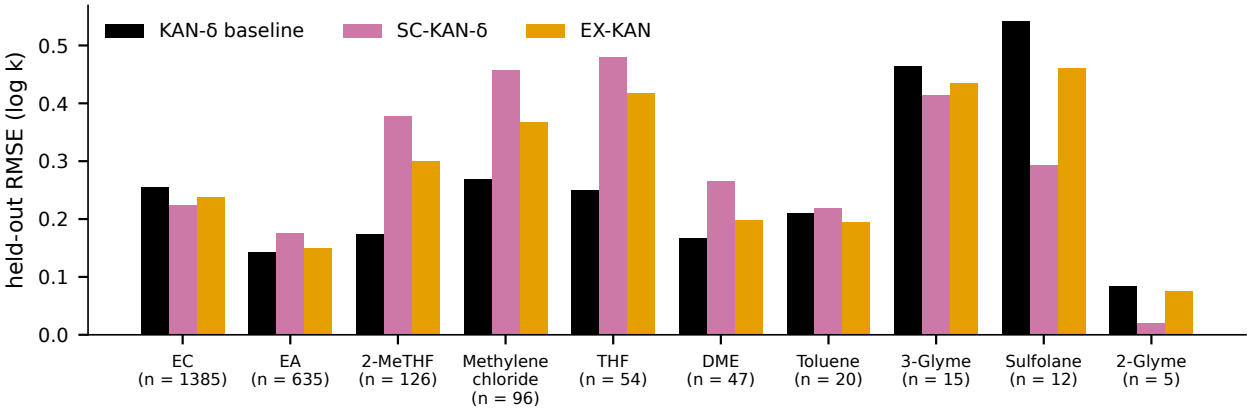


Figure 12: Leave-one-co-solvent-out RMSE per fold (5-seed ensembles, early state) for the KAN-δ baseline, SC-KAN-δ and EX-KAN. Unweighted means: KAN-δ baseline 0.2554; SC-KAN-δ 0.2923; EX-KAN 0.2834. The EA fold uses the fold A rows.

F Supplementary results

These passages were moved from the main text for length; they carry the same evidence and the same labels.

F.1 Freezing the second layer (exploratory arm)

Freezing the second layer instead of the first leaves a non-convex problem. The synthetic models still certified (16 of 16), but the real first layer grew by factors of 10^4 to 10^{10}, mostly through cancellation across its edges, and its endpoints were highly degenerate, with 21–41 near-flat directions; one endpoint that passed the certificate turned out, on a longer restart, to be still drifting outward (exploratory arm; post-hoc follow-up).

F.2 Structured alternatives

Structured alternatives. We tested two architectures designed to remove runaway routes by construction, against the early-stopped baseline, with leave-one-co-solvent-out (LOCO) validation:

Table 17: Structured alternatives against the early-stopped baseline.

Model	Idea	LOCO RMSE	Converged	Seed agreement
Baseline KAN-δ	As above	0.255	0/10	0.112
Shared-curve KAN	One curve per input, mixed by penalised weights	0.292	10/10	0.187 early; 0.085 polished
Two-anchor excess model	Pure-solvent offset plus $x(1-x)$ mixing term	0.283	not tested	not tested

Seed agreement is the median prediction difference between seeds (lower is better). The shared-curve model converged every time, ending with no near-flat directions where the baseline had a median of 106 (range 86–116, at non-stationary points; Figure 5a), and it was less accurate. At the early-stopped state its seeds agreed less than the baseline's (the pre-registered comparison); after polishing to certified minima their median disagreement fell to 0.085, with some seed pairs reaching the same minimum (post-hoc). The excess model's "pure co-solvent offset" tracked viscosity rather than dielectric constant (rank correlation -0.85, post hoc), reflecting that no training row contains a pure co-solvent.

F.3 The regularised teacher and gauge fixing

A known solution becomes recoverable. Under the regularised loss, the original teacher is no longer a minimum. Optimising it under the same loss gives a "regularised teacher" computing nearly the same function (correlation 0.9992). Students match its units at 0.87–0.90 and its edges at 0.97–0.98, against 0.68–0.72 and 0.84–0.91 for the original teacher. This comparison was defined after the results and is exploratory.

The penalty acts partly as gauge fixing. At the certified minima, the layer-2 constant shift, an essentially exact symmetry, owes 96% of its curvature to μQ: the regulariser selects one representative. The layer-1 constant shift (46%) and unit rescaling (24%) are only approximate symmetries in this architecture, because Gaussian bumps cannot represent constants exactly and tanh breaks rescaling.

F.4 What a short path cannot tell

What a short path cannot tell. We also predicted that a short regularisation path ($\gamma = 10^{-3}$, 10^{-5}, 10^{-7}), summarised by $\mathrm{D} = \|\theta^*(10^{-7})\|^2 \,/\, \|\theta^*(10^{-5})\|^2$, would diverge more strongly as ρ grows. It did not (Spearman 0.71 under P0 and 0.14 under P1, against the required 0.8): D was already high at every dose, including $\rho = 0$. A post-hoc check, with its classification rule fixed before computing, explains why. Where the unregularised polish certified, mostly at low ρ, 30 of 32 paths were heading towards that certified finite minimum or its valley, none moved outward past it, and 2 found a better minimum at smaller norm. Along one direction with curvature λ, $D = \big((\lambda + 2\gamma_1)/(\lambda + 2\gamma_2)\big)^2$ with $\gamma_1 = 10^{-5}$ and $\gamma_2 = 10^{-7}$, so a finite minimum whose curvature lies within or near the window $2{\cdot}10^{-7}$ to $2{\cdot}10^{-5}$ keeps D large while the path

is still approaching it; the observed low-ρ medians (4.54 to 15.05) correspond to $\lambda = 6.7{\cdot}10^{-6}$ to $1.7{\cdot}10^{-5}$. A two-decade D therefore cannot separate a distant finite minimum from a receding one. The receding minimum of Section 6 is established at $\rho = 0.95$ and on the real data; at low ρ the minima are finite but distant; at intermediate ρ these tests do not settle the question. A secondary prediction, that the affine share of the parameter change would rise steadily with ρ, failed.

F.5 The exact symmetry of B-spline edges

Because B-splines sum exactly to one, constants can be moved between a unit's incoming edges and its bias without changing the loss: an exact 15-dimensional symmetry in the [4, 3, 1] model, which holds only approximately for RBF edges. It made the solver hang at converged points and would have made every certified point look degenerate. We excluded these analytically constructed directions, never fitted ones, from certification and from the degeneracy count; exact symmetries are redundancy in the parameterisation, while the runaway directions, which are approximate and data-dependent, remain subject to every check.